

Copy_of_Copy_of_Copy_of_5_DonorsChoose_LR

March 19, 2019

1 DonorsChoose

DonorsChoose.org receives hundreds of thousands of project proposals each year for classroom projects in need of funding. Right now, a large number of volunteers is needed to manually screen each submission before it's approved to be posted on the DonorsChoose.org website.

Next year, DonorsChoose.org expects to receive close to 500,000 project proposals. As a result, there are three main p

How to scale current manual processes and resources to screen 500,000 projects so that they can be posted as quickly a

How to increase the consistency of project vetting across different volunteers to improve the experience for teach

How to focus volunteer time on the applications that need the most assistance

The goal of the competition is to predict whether or not a DonorsChoose.org project proposal submitted by a teacher will be approved, using the text of project descriptions as well as additional metadata about the project, teacher, and school. DonorsChoose.org can then use this information to identify projects most likely to need further review before approval.

1.1 About the DonorsChoose Data Set

The train.csv data set provided by DonorsChoose contains the following features:

Feature	Description
project_id	A unique identifier for the proposed project. Example: p036502

project_title | Title of the project. **Examples:**

Art Will Make You Happy!

First Grade Fun

project_grade_category | Grade level of students for which the project is targeted. One of the following enumerated values:

Grades PreK-2

Grades 3-5

Grades 6-8

Grades 9-12

project_subject_categories | One or more (comma-separated) subject categories for the project from the following enumerated list of values:

Applied Learning
 Care & Hunger
 Health & Sports
 History & Civics
 Literacy & Language
 Math & Science
 Music & The Arts
 Special Needs
 Warmth

Examples:

Music & The Arts
 Literacy & Language, Math & Science

school_state | State where school is located ([Two-letter U.S. postal code](#)). **Example:** WY
 project_subject_subcategories | One or more (comma-separated) subject subcategories for the project. **Examples:**

Literacy
 Literature & Writing, Social Sciences

project_resource_summary | An explanation of the resources needed for the project. **Example:**
 My students need hands on literacy materials to manage sensory needs!

project_essay_1 | First application essay

project_essay_2 | *Second application essay* project_essay_3 | Third application essay
 project_essay_4 | *Fourth application essay* project_submitted_datetime | Datetime when project application was submitted. **Example:** 2016-04-28 12:43:56.245

teacher_id | A unique identifier for the teacher of the proposed project. **Example:**
 bdf8baa8fedef6bfec7ae4ff1c15c56

teacher_prefix | Teacher's title. One of the following enumerated values:

nan
 Dr.
 Mr.
 Mrs.
 Ms.
 Teacher.

teacher_number_of_previously_posted_projects | Number of project applications previously submitted by the same teacher. **Example:** 2

* See the section Notes on the Essay Data for more details about these features.

Additionally, the resources.csv data set provides more data about the resources required for each project. Each line in this file represents a resource required by a project:

Feature	Description
id	A project_id value from the train.csv file. Example: p036502
description	Description of the resource. Example: Tenor Saxophone Reeds, Box of 25

Feature	Description
quantity	Quantity of the resource required.
price	Price of the resource required. Example: 9.95

Note: Many projects require multiple resources. The id value corresponds to a `project_id` in `train.csv`, so you use it as a key to retrieve all resources needed for a project:

The data set contains the following label (the value you will attempt to predict):

Label	Description
<code>project_is_approved</code>	Arbitrary flag indicating whether DonorsChoose approved the project. A value of 0 indicates the project was not approved, and a value of 1 indicates the project was approved.

1.1.1 Notes on the Essay Data

Prior to May 17, 2016, the prompts for the essays were as follows:

project_essay_1: "Introduce us to your classroom"

project_essay_2: "Tell us more about your students"

project_essay_3: "Describe how your students will use the materials you're requesting"

project_essay_3: "Close by sharing why your project will make a difference"

Starting on May 17, 2016, the number of essays was reduced from 4 to 2, and the prompts for the first 2 essays were changed to the following:

project_essay_1: "Describe your students: What makes your students special? Specific details about their background, your neighborhood, and your school are all helpful."

project_essay_2: "About your project: How will these materials make a difference in your students' learning and improve their school lives?"

For all projects with `project_submitted_datetime` of 2016-05-17 and later, the values of `project_essay_3` and `project_essay_4` will be NaN.

```
In [1]: %matplotlib inline
import warnings
warnings.filterwarnings("ignore")

import sqlite3
import pandas as pd
```

```

import numpy as np
import nltk
import string
import matplotlib.pyplot as plt
import seaborn as sns
from sklearn.feature_extraction.text import TfidfTransformer
from sklearn.feature_extraction.text import TfidfVectorizer

from sklearn.feature_extraction.text import CountVectorizer
from sklearn.metrics import confusion_matrix
from sklearn import metrics
from sklearn.metrics import roc_curve, auc
from nltk.stem.porter import PorterStemmer

import re
# Tutorial about Python regular expressions: https://pymotw.com/2/re/
import string
from nltk.corpus import stopwords
from nltk.stem import PorterStemmer
from nltk.stem.wordnet import WordNetLemmatizer

from gensim.models import Word2Vec
from gensim.models import KeyedVectors
import pickle

from tqdm import tqdm
import os

from plotly import plotly
import plotly.offline as offline
import plotly.graph_objs as go
offline.init_notebook_mode()
from collections import Counter

```

1.2 1.1 Reading Data

In [2]: #since im using google_colab, i have to mount the gdrive folder for accessing the files

```

from google.colab import drive
drive.mount('/content/gdrive')

```

Drive already mounted at /content/gdrive; to attempt to forcibly remount, call drive.mount("/content/gdrive", force_remount=True)

In [0]: #reading the datasets, i have taken only 5000 datapoints into consideration for avoiding mermory issues

```

project_data = pd.read_csv('/content/gdrive/My Drive/Colab Notebooks/Assignments_DonorsChoose_2019/project_data.csv')
resource_data = pd.read_csv('/content/gdrive/My Drive/Colab Notebooks/Assignments_DonorsChoose_2019/resource_data.csv')

```

```
In [4]: print("Number of data points in train data", project_data.shape)
        print('-'*50)
        print("The attributes of data :", project_data.columns.values)
```

Number of data points in train data (50000, 17)

```
-----
The attributes of data : ['Unnamed: 0' 'id' 'teacher_id' 'teacher_prefix' 'school_state'
'project_submitted_datetime' 'project_grade_category'
'project_subject_categories' 'project_subject_subcategories'
'project_title' 'project_essay_1' 'project_essay_2' 'project_essay_3'
'project_essay_4' 'project_resource_summary'
'teacher_number_of_previously_posted_projects' 'project_is_approved']
```

```
In [5]: # how to replace elements in list python: https://stackoverflow.com/a/2582163/4084039
        cols = ['Date' if x=='project_submitted_datetime' else x for x in list(project_data.columns)]
```

```
#sort dataframe based on time pandas python: https://stackoverflow.com/a/49702492/4084039
project_data['Date'] = pd.to_datetime(project_data['project_submitted_datetime'])
project_data.drop('project_submitted_datetime', axis=1, inplace=True)
project_data.sort_values(by=['Date'], inplace=True)
```

```
# how to reorder columns pandas python: https://stackoverflow.com/a/13148611/4084039
project_data = project_data[cols]
```

```
project_data.head(2)
```

```
Out[5]:      Unnamed: 0      id      teacher_id teacher_prefix \
473      100660 p234804 cbc0e38f522143b86d372f8b43d4cff3      Mrs.
41558      33679 p137682 06f6e62e17de34fcf81020c77549e1d5      Mrs.

      school_state      Date project_grade_category \
473      GA 2016-04-27 00:53:00      Grades PreK-2
41558      WA 2016-04-27 01:05:25      Grades 3-5

      project_subject_categories project_subject_subcategories \
473      Applied Learning      Early Development
41558      Literacy & Language      Literacy

      project_title \
473      Flexible Seating for Flexible Learning
41558      Going Deep: The Art of Inner Thinking!

      project_essay_1 \
473      I recently read an article about giving studen...
```

41558 My students crave challenge, they eat obstacle...

```
project_essay_2 \
473 I teach at a low-income (Title 1) school. Ever...
41558 We are an urban, public k-5 elementary school...
```

```
project_essay_3 \
473 We need a classroom rug that we can use as a c...
41558 With the new common core standards that have b...
```

```
project_essay_4 \
473 Benjamin Franklin once said, \"Tell me and I f...
41558 These remarkable gifts will provide students w...
```

```
project_resource_summary \
473 My students need flexible seating in the class...
41558 My students need copies of the New York Times ...
```

```
teacher_number_of_previously_posted_projects project_is_approved
473 2 1
41558 2 1
```

```
In [6]: print("Number of data points in train data", resource_data.shape)
print(resource_data.columns.values)
resource_data.head(2)
```

Number of data points in train data (1541272, 4)
['id' 'description' 'quantity' 'price']

```
Out[6]:      id      description  quantity \
0 p233245 LC652 - Lakeshore Double-Space Mobile Drying Rack      1
1 p069063 Bouncy Bands for Desks (Blue support pipes)      3

      price
0 149.00
1 14.95
```

1.3 1.2 preprocessing of project_subject_categories

```
In [0]: catogories = list(project_data['project_subject_categories'].values)
# remove special characters from list of strings python: https://stackoverflow.com/a/47301924/4084039

# https://www.geeksforgeeks.org/removing-stop-words-nltk-python/
# https://stackoverflow.com/questions/23669024/how-to-strip-a-specific-word-from-a-string
# https://stackoverflow.com/questions/8270092/remove-all-whitespace-in-a-string-in-python
cat_list = []
for i in catogories:
    temp = ""
```

```

# consider we have text like this "Math & Science, Warmth, Care & Hunger"
for j in i.split(','): # it will split it in three parts ["Math & Science", "Warmth", "Care & Hunger"]
    if 'The' in j.split(): # this will split each of the category based on space "Math & Science"=> "Math
        j=j.replace('The','') # if we have the words "The" we are going to replace it with ""(i.e removing '
    j = j.replace(' ','') # we are placing all the ' '(space) with ''(empty) ex:"Math & Science"=>"Math&
    temp+=j.strip()+" " #" abc ".strip() will return "abc", remove the trailing spaces
    temp = temp.replace('&','_') # we are replacing the & value into
cat_list.append(temp.strip())

project_data['clean_categories'] = cat_list
project_data.drop(['project_subject_categories'], axis=1, inplace=True)

from collections import Counter
my_counter = Counter()
for word in project_data['clean_categories'].values:
    my_counter.update(word.split())

cat_dict = dict(my_counter)
sorted_cat_dict = dict(sorted(cat_dict.items(), key=lambda kv: kv[1]))

```

1.4 1.3 preprocessing of project_subject_subcategories

```

In [0]: sub_catogories = list(project_data['project_subject_subcategories'].values)
# remove special characters from list of strings python: https://stackoverflow.com/a/47301924/4084039

# https://www.geeksforgeeks.org/removing-stop-words-nltk-python/
# https://stackoverflow.com/questions/23669024/how-to-strip-a-specific-word-from-a-string
# https://stackoverflow.com/questions/8270092/remove-all-whitespace-in-a-string-in-python

sub_cat_list = []
for i in sub_catogories:
    temp = ""
    # consider we have text like this "Math & Science, Warmth, Care & Hunger"
    for j in i.split(','): # it will split it in three parts ["Math & Science", "Warmth", "Care & Hunger"]
        if 'The' in j.split(): # this will split each of the category based on space "Math & Science"=> "Math
            j=j.replace('The','') # if we have the words "The" we are going to replace it with ""(i.e removing '
        j = j.replace(' ','') # we are placing all the ' '(space) with ''(empty) ex:"Math & Science"=>"Math&
        temp +=j.strip()+" " #" abc ".strip() will return "abc", remove the trailing spaces
        temp = temp.replace('&','_')
    sub_cat_list.append(temp.strip())

project_data['clean_subcategories'] = sub_cat_list
project_data.drop(['project_subject_subcategories'], axis=1, inplace=True)

# count of all the words in corpus python: https://stackoverflow.com/a/22898595/4084039
my_counter = Counter()
for word in project_data['clean_subcategories'].values:
    my_counter.update(word.split())

```

```
sub_cat_dict = dict(my_counter)
sorted_sub_cat_dict = dict(sorted(sub_cat_dict.items(), key=lambda kv: kv[1]))
```

1.5 1.3 Text preprocessing

In [0]: # merge two column text dataframe:

```
project_data["essay"] = project_data["project_essay_1"].map(str) + \
    project_data["project_essay_2"].map(str) + \
    project_data["project_essay_3"].map(str) + \
    project_data["project_essay_4"].map(str)
```

In [10]: project_data.head(2)

```
Out[10]:
```

	Unnamed: 0	id	teacher_id	teacher_prefix	\
473	100660	p234804	cbc0e38f522143b86d372f8b43d4cff3	Mrs.	
41558	33679	p137682	06f6e62e17de34fcf81020c77549e1d5	Mrs.	

	school_state	Date	project_grade_category	\
473	GA	2016-04-27 00:53:00	Grades PreK-2	
41558	WA	2016-04-27 01:05:25	Grades 3-5	

	project_title	\
473	Flexible Seating for Flexible Learning	
41558	Going Deep: The Art of Inner Thinking!	

	project_essay_1	\
473	I recently read an article about giving studen...	
41558	My students crave challenge, they eat obstacle...	

	project_essay_2	\
473	I teach at a low-income (Title 1) school. Ever...	
41558	We are an urban, public k-5 elementary school...	

	project_essay_3	\
473	We need a classroom rug that we can use as a c...	
41558	With the new common core standards that have b...	

	project_essay_4	\
473	Benjamin Franklin once said, \"Tell me and I f...	
41558	These remarkable gifts will provide students w...	

	project_resource_summary	\
473	My students need flexible seating in the class...	
41558	My students need copies of the New York Times ...	

	teacher_number_of_previously_posted_projects	project_is_approved	\
473	2	1	

41558

2

1

```

clean_categories clean_subcategories \
473   AppliedLearning   EarlyDevelopment
41558 Literacy_Language           Literacy

```

```

essay
473   I recently read an article about giving studen...
41558 My students crave challenge, they eat obstacle...

```

In [0]: ##### 1.4.2.3 Using Pretrained Models: TFIDF weighted W2V

```

In [12]: # printing some random reviews
print(project_data['essay'].values[0])
print("="*50)
print(project_data['essay'].values[150])
print("="*50)
print(project_data['essay'].values[1000])
print("="*50)
print(project_data['essay'].values[20000])
print("="*50)
#print(project_data['essay'].values[99999])
#print("="*50)

```

I recently read an article about giving students a choice about how they learn. We already set goals; why not let t
=====

At the beginning of every class we start out with a Math Application problem to help students see the relevance o
=====

My students love coming to school and they love learning. I strive daily to make our classroom a relaxed, comfort
=====

I teach at a Title 1 school, with 73% of my students who receive free/reduced lunch. Our school provides free bre
=====

In [0]: # <https://stackoverflow.com/a/47091490/4084039>

```

import re

def decontracted(phrase):
    # specific
    phrase = re.sub(r"won't", "will not", phrase)
    phrase = re.sub(r"can't", "can not", phrase)

    # general
    phrase = re.sub(r"n't", " not", phrase)
    phrase = re.sub(r"\re", " are", phrase)
    phrase = re.sub(r"\s", " is", phrase)
    phrase = re.sub(r"\d", " would", phrase)
    phrase = re.sub(r"\ll", " will", phrase)
    phrase = re.sub(r"\t", " not", phrase)

```

```

phrase = re.sub(r"\ve", " have", phrase)
phrase = re.sub(r"\m", " am", phrase)
return phrase

```

```

In [14]: sent = decontracted(project_data['essay'].values[20000])
print(sent)
print("="*50)

```

I teach at a Title 1 school, with 73% of my students who receive free/reduced lunch. Our school provides free bre
=====

```

In [15]: # \r \n \t remove from string python: http://texthandler.com/info/remove-line-breaks-python/
sent = sent.replace('\r', ' ')
sent = sent.replace('\n', ' ')
sent = sent.replace('\t', ' ')
print(sent)

```

I teach at a Title 1 school, with 73% of my students who receive free/reduced lunch. Our school provides free bre

```

In [16]: #remove spacial character: https://stackoverflow.com/a/5843547/4084039
sent = re.sub('[^A-Za-z0-9]+', ' ', sent)
print(sent)

```

I teach at a Title 1 school with 73 of my students who receive free reduced lunch Our school provides free breakfa

```

In [0]: # https://gist.github.com/sebleier/554280
# we are removing the words from the stop words list: 'no', 'nor', 'not'
stopwords= ['i', 'me', 'my', 'myself', 'we', 'our', 'ours', 'ourselves', 'you', "you're", "you've", \
            "you'll", "you'd", 'your', 'yours', 'yourself', 'yourselves', 'he', 'him', 'his', 'himself', \
            'she', "she's", 'her', 'hers', 'herself', 'it', "it's", 'its', 'itself', 'they', 'them', 'their', \
            'theirs', 'themselves', 'what', 'which', 'who', 'whom', 'this', 'that', "that'll", 'these', 'those', \
            'am', 'is', 'are', 'was', 'were', 'be', 'been', 'being', 'have', 'has', 'had', 'having', 'do', 'does', \
            'did', 'doing', 'a', 'an', 'the', 'and', 'but', 'if', 'or', 'because', 'as', 'until', 'while', 'of', \
            'at', 'by', 'for', 'with', 'about', 'against', 'between', 'into', 'through', 'during', 'before', 'after', \
            'above', 'below', 'to', 'from', 'up', 'down', 'in', 'out', 'on', 'off', 'over', 'under', 'again', 'further', \
            'then', 'once', 'here', 'there', 'when', 'where', 'why', 'how', 'all', 'any', 'both', 'each', 'few', 'more', \
            'most', 'other', 'some', 'such', 'only', 'own', 'same', 'so', 'than', 'too', 'very', \
            's', 't', 'can', 'will', 'just', 'don', "don't", 'should', "should've", 'now', 'd', 'll', 'm', 'o', 're', \
            've', 'y', 'ain', 'aren', "aren't", 'couldn', "couldn't", 'didn', "didn't", 'doesn', "doesn't", 'hadn', \
            "hadn't", 'hasn', "hasn't", 'haven', "haven't", 'isn', "isn't", 'ma', 'mightn', "mightn't", 'mustn', \
            "mustn't", 'needn', "needn't", 'shan', "shan't", 'shouldn', "shouldn't", 'wasn', "wasn't", 'weren', \
            'won', "won't", 'wouldn', "wouldn't"]

```

```

In [18]: # Combining all the above stundents
from tqdm import tqdm
preprocessed_essays = []

```

```
# tqdm is for printing the status bar
for sentence in tqdm(project_data['essay'].values):
    sent = decontracted(sentence)
    sent = sent.replace('\r', ' ')
    sent = sent.replace('\n', ' ')
    sent = sent.replace('\n', ' ')
    sent = re.sub('[^A-Za-z0-9]+', ' ', sent)
    # https://gist.github.com/sebleier/554280
    sent = ' '.join(e for e in sent.split() if e.lower() not in stopwords)
    preprocessed_essays.append(sent.lower().strip())
```

100%| 50000/50000 [00:29<00:00, 1667.89it/s]

In [19]: # after preprocessing
preprocessed_essays[20000]

Out[19]: 'teach title 1 school 73 students receive free reduced lunch school provides free breakfast students special

1.4 Preprocessing of project_title

In [20]: # similarly you can preprocess the titles also

```
from tqdm import tqdm
preprocessed_project_title = []
# tqdm is for printing the status bar
for sentence in tqdm(project_data['project_title'].values):
    sent = decontracted(sentence)
    sent = sent.replace('\r', ' ')
    sent = sent.replace('\n', ' ')
    sent = sent.replace('\n', ' ')
    sent = re.sub('[^A-Za-z0-9]+', ' ', sent)
    # https://gist.github.com/sebleier/554280
    sent = ' '.join(e for e in sent.split() if e not in stopwords)
    preprocessed_project_title.append(sent.lower().strip())
```

100%| 50000/50000 [00:01<00:00, 35329.06it/s]

1.6 1.5 Preparing data for models

In [21]: project_data.columns

Out[21]: Index(['Unnamed: 0', 'id', 'teacher_id', 'teacher_prefix', 'school_state',
'Date', 'project_grade_category', 'project_title', 'project_essay_1',
'project_essay_2', 'project_essay_3', 'project_essay_4',
'project_resource_summary',
'teacher_number_of_previously_posted_projects', 'project_is_approved',
'clean_categories', 'clean_subcategories', 'essay'],
dtype='object')

we are going to consider

- school_state : categorical data
- clean_categories : categorical data
- clean_subcategories : categorical data
- project_grade_category : categorical data
- teacher_prefix : categorical data
- project_title : text data
- text : text data
- project_resource_summary: text data (optinal)
- quantity : numerical (optinal)
- teacher_number_of_previously_posted_projects : numerical
- price : numerical

1.6.1 Modifying DataSet (essay & project_title)

```
In [0]: project_data['clean_essay'] = preprocessed_essays
        project_data['clean_project_title'] = preprocessed_project_title
        project_data.drop(['essay'], axis=1, inplace=True)
        project_data.drop(['project_title'], axis=1, inplace=True)
```

```
In [23]: project_data.head(1)
```

```
Out[23]:   Unnamed: 0      id      teacher_id teacher_prefix \
473      100660  p234804  cbc0e38f522143b86d372f8b43d4cff3  Mrs.

      school_state      Date project_grade_category \
473      GA 2016-04-27 00:53:00      Grades PreK-2

      project_essay_1 \
473  I recently read an article about giving studen...

      project_essay_2 \
473  I teach at a low-income (Title 1) school. Ever...

      project_essay_3 \
473  We need a classroom rug that we can use as a c...

      project_essay_4 \
473  Benjamin Franklin once said, "Tell me and I f...

      project_resource_summary \
473  My students need flexible seating in the class...

      teacher_number_of_previously_posted_projects  project_is_approved \
473      2      1
```

```

clean_categories clean_subcategories \
473 AppliedLearning EarlyDevelopment

```

```

clean_essay \
473 recently read article giving students choice l...

```

```

clean_project_title
473 flexible seating flexible learning

```

```

In [24]: y = project_data['project_is_approved'].values
project_data.drop(['project_is_approved'], axis=1, inplace=True)
project_data.head(1)

```

```

Out[24]: Unnamed: 0 id teacher_id teacher_prefix \
473 100660 p234804 cbc0e38f522143b86d372f8b43d4cff3 Mrs.

```

```

school_state Date project_grade_category \
473 GA 2016-04-27 00:53:00 Grades PreK-2

```

```

project_essay_1 \
473 I recently read an article about giving studen...

```

```

project_essay_2 \
473 I teach at a low-income (Title 1) school. Ever...

```

```

project_essay_3 \
473 We need a classroom rug that we can use as a c...

```

```

project_essay_4 \
473 Benjamin Franklin once said, "Tell me and I f...

```

```

project_resource_summary \
473 My students need flexible seating in the class...

```

```

teacher_number_of_previously_posted_projects clean_categories \
473 2 AppliedLearning

```

```

clean_subcategories clean_essay \
473 EarlyDevelopment recently read article giving students choice l...

```

```

clean_project_title
473 flexible seating flexible learning

```

```

In [0]: X = project_data

```

2.1 Splitting data into Train and cross validation(or test): Stratified Sampling

```

In [0]: # please write all the code with proper documentation, and proper titles for each subsection
# go through documentations and blogs before you start coding

```

```

# first figure out what to do, and then think about how to do.
# reading and understanding error messages will be very much helpfull in debugging your code
# when you plot any graph make sure you use
# a. Title, that describes your plot, this will be very helpful to the reader
# b. Legends if needed
# c. X-axis label
# d. Y-axis label

# train test split
from sklearn.model_selection import train_test_split
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.33, stratify=y)
#X_train, X_cv, y_train, y_cv = train_test_split(X_train, y_train, test_size=0.33, stratify=y_train)

```

In [27]: X_train.head(2)

```

Out[27]:      Unnamed: 0      id      teacher_id teacher_prefix \
8971      36401  p111136  a703366918b87aaa7eaf013321ccd98f      Mrs.
26908      145729  p228634  5c250f8b1523438faa7b40e59b21df0a      Mrs.

      school_state      Date project_grade_category \
8971      GA 2016-08-22 12:53:43      Grades 6-8
26908      KY 2016-08-02 17:32:08      Grades 3-5

      project_essay_1 \
8971  We are a Montessori school with a deep history...
26908  My classroom, of nearly 30 students, is full o...

      project_essay_2 project_essay_3 \
8971  We are moving to a healthier student body and ...      NaN
26908  When I first started thinking about the idea o...      NaN

      project_essay_4      project_resource_summary \
8971      NaN  My students need some fitness stuff.  We are m...
26908      NaN  My students need active seating and standing o...

      teacher_number_of_previously_posted_projects clean_categories \
8971      5      Health_Sports
26908      4      Health_Sports

      clean_subcategories      clean_essay \
8971      Gym_Fitness  montessori school deep history student centere...
26908      Gym_Fitness  classroom nearly 30 students full students wan...

      clean_project_title
8971      fitness frenzy
26908      moving success

```

In [28]: resource_data.head(1)

```
Out[28]:      id                description quantity price
0  p233245  LC652 - Lakeshore Double-Space Mobile Drying Rack      1  149.0
```

2.2 Make Data Model Ready: encoding numerical, categorical features

1.6.2 clean_categories

```
In [29]: vectorizer = CountVectorizer()
        vectorizer.fit(X_train['clean_categories'].values) # fit has to happen only on train data

        # we use the fitted CountVectorizer to convert the text to vector
        X_train_cc_ohe = vectorizer.transform(X_train['clean_categories'].values)
        #X_cv_cc_ohe = vectorizer.transform(X_cv['clean_categories'].values)
        X_test_cc_ohe = vectorizer.transform(X_test['clean_categories'].values)

        print("After vectorizations")
        print(X_train_cc_ohe.shape, y_train.shape)
        #print(X_cv_cc_ohe.shape, y_cv.shape)
        print(X_test_cc_ohe.shape, y_test.shape)
        print(vectorizer.get_feature_names())
        print("="*100)
```

After vectorizations

(33500, 9) (33500,)

(16500, 9) (16500,)

['appliedlearning', 'care_hunger', 'health_sports', 'history_civics', 'literacy_language', 'math_science', 'music_a

=====

1.6.3 clean_subcategories

```
In [30]: vectorizer = CountVectorizer()
        vectorizer.fit(X_train['clean_subcategories'].values) # fit has to happen only on train data

        # we use the fitted CountVectorizer to convert the text to vector
        X_train_csc_ohe = vectorizer.transform(X_train['clean_subcategories'].values)
        #X_cv_csc_ohe = vectorizer.transform(X_cv['clean_subcategories'].values)
        X_test_csc_ohe = vectorizer.transform(X_test['clean_subcategories'].values)

        print("After vectorizations")
        print(X_train_csc_ohe.shape, y_train.shape)
        #print(X_cv_csc_ohe.shape, y_cv.shape)
        print(X_test_csc_ohe.shape, y_test.shape)
        print(vectorizer.get_feature_names())
        print("="*100)
```

After vectorizations

(33500, 30) (33500,)

(16500, 30) (16500,)

```
['appliedsciences', 'care_hunger', 'charactereducation', 'civics_government', 'college_careerprep', 'communityserv  
=====
```

1.6.4 school_state

```
In [31]: vectorizer = CountVectorizer()  
        vectorizer.fit(X_train['school_state'].values) # fit has to happen only on train data  
  
# we use the fitted CountVectorizer to convert the text to vector  
X_train_state_ohe = vectorizer.transform(X_train['school_state'].values)  
#X_cv_state_ohe = vectorizer.transform(X_cv['school_state'].values)  
X_test_state_ohe = vectorizer.transform(X_test['school_state'].values)  
  
print("After vectorizations")  
print(X_train_state_ohe.shape, y_train.shape)  
#print(X_cv_state_ohe.shape, y_cv.shape)  
print(X_test_state_ohe.shape, y_test.shape)  
print(vectorizer.get_feature_names())  
print("="*100)
```

After vectorizations

```
(33500, 51) (33500,)
```

```
(16500, 51) (16500,)
```

```
['ak', 'al', 'ar', 'az', 'ca', 'co', 'ct', 'dc', 'de', 'fl', 'ga', 'hi', 'ia', 'id', 'il', 'in', 'ks', 'ky', 'la', 'ma', 'md', 'me', 'mi', 'm  
=====
```

1.6.5 teacher_prefix

```
In [32]: vectorizer = CountVectorizer()  
        vectorizer.fit(X_train['teacher_prefix'].values.astype('U')) # fit has to happen only on train data  
  
# we use the fitted CountVectorizer to convert the text to vector  
X_train_teacher_ohe = vectorizer.transform(X_train['teacher_prefix'].values.astype('U'))  
#X_cv_teacher_ohe = vectorizer.transform(X_cv['teacher_prefix'].values.astype('U'))  
X_test_teacher_ohe = vectorizer.transform(X_test['teacher_prefix'].values.astype('U'))  
  
print("After vectorizations")  
print(X_train_teacher_ohe.shape, y_train.shape)  
#print(X_cv_teacher_ohe.shape, y_cv.shape)  
print(X_test_teacher_ohe.shape, y_test.shape)  
print(vectorizer.get_feature_names())  
print("="*100)
```

After vectorizations

```
(33500, 5) (33500,)
```

```
(16500, 5) (16500,)
```

```
['dr', 'mr', 'mrs', 'ms', 'teacher']
```


1.6.6 project_grade_category

```
=====

In [33]: vectorizer = CountVectorizer()
         vectorizer.fit(X_train['project_grade_category'].values) # fit has to happen only on train data

         # we use the fitted CountVectorizer to convert the text to vector
         X_train_grade_ohe = vectorizer.transform(X_train['project_grade_category'].values)
         #X_cv_grade_ohe = vectorizer.transform(X_cv['project_grade_category'].values)
         X_test_grade_ohe = vectorizer.transform(X_test['project_grade_category'].values)

         print("After vectorizations")
         print(X_train_grade_ohe.shape, y_train.shape)
         #print(X_cv_grade_ohe.shape, y_cv.shape)
         print(X_test_grade_ohe.shape, y_test.shape)
         print(vectorizer.get_feature_names())
         print("="*100)
```

After vectorizations

(33500, 3) (33500,)

(16500, 3) (16500,)

['12', 'grades', 'prek']

2.3 Make Data Model Ready: encoding eassay, and project_title

1.5.2.1 Bag of words

1.6.7 essays

```
In [34]: from sklearn.feature_extraction.text import CountVectorizer
         vectorizer = CountVectorizer(min_df=10,ngram_range=(2,2), max_features=5000)
         vectorizer.fit(X_train['clean_essay'].values) # fit has to happen only on train data

         # we use the fitted CountVectorizer to convert the text to vector
         X_train_essay_bow = vectorizer.transform(X_train['clean_essay'].values)
         #X_cv_essay_bow = vectorizer.transform(X_cv['clean_essay'].values)
         X_test_essay_bow = vectorizer.transform(X_test['clean_essay'].values)

         print("After vectorizations")
         print(X_train_essay_bow.shape, y_train.shape)
         #print(X_cv_essay_bow.shape, y_cv.shape)
         print(X_test_essay_bow.shape, y_test.shape)
         print("="*100)
```

After vectorizations
(33500, 5000) (33500,)
(16500, 5000) (16500,)

1.6.8 project_title

```
In [35]: from sklearn.feature_extraction.text import CountVectorizer
vectorizer = CountVectorizer(min_df=10)
vectorizer.fit(X_train['clean_project_title'].values) # fit has to happen only on train data

# we use the fitted CountVectorizer to convert the text to vector
X_train_pt_bow = vectorizer.transform(X_train['clean_project_title'].values)
#X_cv_pt_bow = vectorizer.transform(X_cv['clean_project_title'].values)
X_test_pt_bow = vectorizer.transform(X_test['clean_project_title'].values)

print("After vectorizations")
print(X_train_pt_bow.shape, y_train.shape)
#print(X_cv_pt_bow.shape, y_cv.shape)
print(X_test_pt_bow.shape, y_test.shape)
print("="*100)
```

After vectorizations
(33500, 1669) (33500,)
(16500, 1669) (16500,)

1.5.2.2 TFIDF vectorizer

1.6.9 essays

```
In [36]: from sklearn.feature_extraction.text import TfidfVectorizer
vectorizer = TfidfVectorizer(min_df=10, ngram_range=(2,2), max_features=5000)
X_train_essay_tfidf = vectorizer.fit_transform(X_train['clean_essay'].values)
#X_cv_essay_tfidf = vectorizer.transform(X_cv['clean_essay'].values)
X_test_essay_tfidf = vectorizer.transform(X_test['clean_essay'].values)

print("Shape of matrix after one hot encodig ",X_train_essay_tfidf.shape)
#print("Shape of matrix after one hot encodig ",X_cv_essay_tfidf.shape)
print("Shape of matrix after one hot encodig ",X_test_essay_tfidf.shape)
```

Shape of matrix after one hot encodig (33500, 5000)
Shape of matrix after one hot encodig (16500, 5000)

1.6.10 project_title

```
In [37]: from sklearn.feature_extraction.text import TfidfVectorizer
vectorizer = TfidfVectorizer(min_df=10)
X_train_pt_tfidf = vectorizer.fit_transform(X_train['clean_project_title'].values)
#X_cv_pt_tfidf = vectorizer.transform(X_cv['clean_project_title'].values)
X_test_pt_tfidf = vectorizer.transform(X_test['clean_project_title'].values)
```

```
print("Shape of matrix after one hot encoding ",X_train_pt_tfidf.shape)
#print("Shape of matrix after one hot encoding ",X_cv_pt_tfidf.shape)
print("Shape of matrix after one hot encoding ",X_test_pt_tfidf.shape)
```

Shape of matrix after one hot encoding (33500, 1669)

Shape of matrix after one hot encoding (16500, 1669)

1.5.2.3 Using Pretrained Models: Avg W2V

1.6.11 essays

Train

```
In [0]: i=0
list_of_sentenceTrain=[]
for sentence in X_train['clean_essay']:
    list_of_sentenceTrain.append(sentence.split())

In [39]: is_your_ram_gt_16g=False
want_to_use_google_w2v = False
want_to_train_w2v = True

if want_to_train_w2v:
    # min_count = 5 considers only words that occurred atleast 5 times
    w2v_model=Word2Vec(list_of_sentenceTrain,min_count=5,size=50, workers=4)
    print(w2v_model.wv.most_similar('great'))
    print('='*50)
    #print(w2v_model.wv.most_similar('worst'))

elif want_to_use_google_w2v and is_your_ram_gt_16g:
    if os.path.isfile('GoogleNews-vectors-negative300.bin'):
        w2v_model=KeyedVectors.load_word2vec_format('GoogleNews-vectors-negative300.bin', binary=True)
        #print(w2v_model.wv.most_similar('great'))
        #print(w2v_model.wv.most_similar('worst'))
    else:
        print("you don't have gogole's word2vec file, keep want_to_train_w2v = True, to train your own w
```

```
[('amazing', 0.731677770614624), ('wonderful', 0.6918390393257141), ('awesome', 0.6692880988121033), ('excellent', 0.6692880988121033)]
=====
```

```
In [40]: w2v_words = list(w2v_model.wv.vocab)
         print("number of words that occurred minimum 5 times ", len(w2v_words))
         print("sample words ", w2v_words[0:50])
```

number of words that occurred minimum 5 times 14633

sample words ['montessori', 'school', 'deep', 'history', 'student', 'centered', 'learning', 'wide', 'range', 'students', 'l

```
In [41]: sent_vectorsPPE_train = []; # the avg-w2v for each sentence/review is stored in this list
         for sent in tqdm(list_of_sentenceTrain): # for each review/sentence
             sent_vec = np.zeros(50) # as word vectors are of zero length 50, you might need to change this to 300
             cnt_words = 0; # num of words with a valid vector in the sentence/review
             for word in sent: # for each word in a review/sentence
                 if word in w2v_words:
                     vec = w2v_model.wv[word]
                     sent_vec += vec
                     cnt_words += 1
             if cnt_words != 0:
                 sent_vec /= cnt_words
             sent_vectorsPPE_train.append(sent_vec)
         print(len(sent_vectorsPPE_train))
         print(len(sent_vectorsPPE_train[0]))
```

100%|| 33500/33500 [02:44<00:00, 203.03it/s]

33500

50

Test

```
In [0]: i=0
         list_of_sentenceTest=[]
         for sentence in X_test['clean_essay']:
             list_of_sentenceTest.append(sentence.split())
```

```
In [43]: is_your_ram_gt_16g=False
         want_to_use_google_w2v = False
         want_to_train_w2v = True

         if want_to_train_w2v:
             # min_count = 5 considers only words that occurred atleast 5 times
```


1.6.12 project_title

Train

In [0]: # Similarly you can vectorize for title also

```
# Train your own Word2Vec model using your own text corpus
i=0
list_of_sentencePTtrain=[]
for sentence in X_train['clean_project_title']:
    list_of_sentencePTtrain.append(sentence.split())
```

```
In [47]: is_your_ram_gt_16g=False
        want_to_use_google_w2v = False
        want_to_train_w2v = True
```

```
try :
```

```
if want_to_train_w2v:
    # min_count = 5 considers only words that occurred atleast 5 times
    w2v_model=Word2Vec(list_of_sentencePTtrain,min_count=5,size=50, workers=4)
    print(w2v_model.wv.most_similar('great'))
    print('='*50)
    print(w2v_model.wv.most_similar('worst'))
```

```
elif want_to_use_google_w2v and is_your_ram_gt_16g:
    if os.path.isfile('GoogleNews-vectors-negative300.bin'):
        w2v_model=KeyedVectors.load_word2vec_format('GoogleNews-vectors-negative300.bin', binary=True)
        print(w2v_model.wv.most_similar('great'))
        print(w2v_model.wv.most_similar('worst'))
    else:
        print("you don't have google's word2vec file, keep want_to_train_w2v = True, to train your own w2v")
```

```
except KeyError :
```

```
    pass
```

```
finally :
```

```
    print("Execution Done")
```

```
[('comic', 0.9938017725944519), ('boys', 0.9927322268486023), ('level', 0.9919877648353577), ('avid', 0.9918404817
```

```
=====
```

Execution Done

```
In [48]: w2v_words = list(w2v_model.wv.vocab)
print("number of words that occurred minimum 5 times ", len(w2v_words))
print("sample words ", w2v_words[0:50])
```

number of words that occurred minimum 5 times 2715

sample words ['fitness', 'frenzy', 'moving', 'success', 'technology', 'third', 'graders', 'creating', 'lifelong', 'readers',

```
In [49]: sent_vectorsPT_train = []; # the avg-w2v for each sentence/review is stored in this list
for sent in tqdm(list_of_sentencePTtrain): # for each review/sentence
    sent_vec = np.zeros(50) # as word vectors are of zero length 50, you might need to change this to 300
    cnt_words = 0; # num of words with a valid vector in the sentence/review
    for word in sent: # for each word in a review/sentence
        if word in w2v_words:
            vec = w2v_model.wv[word]
            sent_vec += vec
            cnt_words += 1
    if cnt_words != 0:
        sent_vec /= cnt_words
    sent_vectorsPT_train.append(sent_vec)
print(len(sent_vectorsPT_train))
print(len(sent_vectorsPT_train[0]))
```

100%| 33500/33500 [00:03<00:00, 10099.97it/s]

33500

50

Test

```
In [0]: i=0
list_of_sentencePT_test=[]
for sentence in X_test['clean_project_title']:
    list_of_sentencePT_test.append(sentence.split())
```

```
In [51]: # Using Google News Word2Vectors
```

```
# in this project we are using a pretrained model by google
# its 3.3G file, once you load this into your memory
# it occupies ~9Gb, so please do this step only if you have >12G of ram
# we will provide a pickle file which contains a dict ,
# and it contains all our corpus words as keys and model[word] as values
# To use this code-snippet, download "GoogleNews-vectors-negative300.bin"
# from https://drive.google.com/file/d/0B7XkCwpI5KDYNlNUTTlSS21pQmM/edit
# it's 1.9GB in size.
```



```

        sent_vec /= cnt_words
        sent_vectorsPT_test.append(sent_vec)
print(len(sent_vectorsPT_test))
print(len(sent_vectorsPT_test[0]))

```

```
100%| 16500/16500 [00:01<00:00, 12716.12it/s]
```

```
16500
```

```
50
```

1.5.2.3 Using Pretrained Models: TFIDF weighted W2V

1.6.13 essays

Train

```

In [0]: # S = ["abc def pqr", "def def def abc", "pqr pqr def"]
        model = TfidfVectorizer()
        model.fit(X_train['clean_essay'])
        # we are converting a dictionary with word as a key, and the idf as a value
        dictionary = dict(zip(model.get_feature_names(), list(model.idf_)))

In [55]: # TF-IDF weighted Word2Vec
        tfidf_feat = model.get_feature_names() # tfidf words/col-names
        # final_tf_idf is the sparse matrix with row= sentence, col=word and cell_val = tfidf

        tfidf_sent_vectors_essay_train = []; # the tfidf-w2v for each sentence/review is stored in this list
        row=0;
        for sent in tqdm(list_of_sentenceTrain): # for each review/sentence
            sent_vec = np.zeros(50) # as word vectors are of zero length
            weight_sum = 0; # num of words with a valid vector in the sentence/review
            for word in sent: # for each word in a review/sentence
                if word in w2v_words and word in tfidf_feat:
                    vec = w2v_model.wv[word]
                    # tf_idf = tf_idf_matrix[row, tfidf_feat.index(word)]
                    # to reduce the computation we are
                    # dictionary[word] = idf value of word in whole corpus
                    # sent.count(word) = tf value of word in this review
                    tf_idf = dictionary[word]*(sent.count(word)/len(sent))
                    sent_vec += (vec * tf_idf)
                    weight_sum += tf_idf
            if weight_sum != 0:
                sent_vec /= weight_sum
            tfidf_sent_vectors_essay_train.append(sent_vec)
            row += 1

```

100%|| 33500/33500 [27:28<00:00, 20.32it/s]

Test

```
In [0]: # S = ["abc def pqr", "def def def abc", "pqr pqr def"]
        model = TfidfVectorizer()
        model.fit(X_test['clean_essay'])
        # we are converting a dictionary with word as a key, and the idf as a value
        dictionary = dict(zip(model.get_feature_names(), list(model.idf_)))

In [57]: # TF-IDF weighted Word2Vec
        tfidf_feat = model.get_feature_names() # tfidf words/col-names
        # final_tf_idf is the sparse matrix with row= sentence, col=word and cell_val = tfidf

        tfidf_sent_vectors_essay_test = []; # the tfidf-w2v for each sentence/review is stored in this list
        row=0;
        for sent in tqdm(list_of_sentenceTest): # for each review/sentence
            sent_vec = np.zeros(50) # as word vectors are of zero length
            weight_sum = 0; # num of words with a valid vector in the sentence/review
            for word in sent: # for each word in a review/sentence
                if word in w2v_words and word in tfidf_feat:
                    vec = w2v_model.wv[word]
                    #
                    tf_idf = tf_idf_matrix[row, tfidf_feat.index(word)]
                    # to reduce the computation we are
                    # dictionary[word] = idf value of word in whole corpus
                    # sent.count(word) = tf value of word in this review
                    tf_idf = dictionary[word]*(sent.count(word)/len(sent))
                    sent_vec += (vec * tf_idf)
                    weight_sum += tf_idf
            if weight_sum != 0:
                sent_vec /= weight_sum
            tfidf_sent_vectors_essay_test.append(sent_vec)
            row += 1
```

100%|| 16500/16500 [10:44<00:00, 25.60it/s]

1.6.14 project_title

Train

```
In [0]: # S = ["abc def pqr", "def def def abc", "pqr pqr def"]
        model = TfidfVectorizer()
        model.fit(X_train['clean_project_title'])
        # we are converting a dictionary with word as a key, and the idf as a value
        dictionary = dict(zip(model.get_feature_names(), list(model.idf_)))
```

```

In [59]: # TF-IDF weighted Word2Vec
tfidf_feat = model.get_feature_names() # tfidf words/col-names
# final_tf_idf is the sparse matrix with row= sentence, col=word and cell_val = tfidf

tfidf_sent_vectorsPT_train = []; # the tfidf-w2v for each sentence/review is stored in this list
row=0;
for sent in tqdm(list_of_sentencePTtrain): # for each review/sentence
    sent_vec = np.zeros(50) # as word vectors are of zero length
    weight_sum =0; # num of words with a valid vector in the sentence/review
    for word in sent: # for each word in a review/sentence
        if word in w2v_words and word in tfidf_feat:
            vec = w2v_model.wv[word]
#            tf_idf = tf_idf_matrix[row, tfidf_feat.index(word)]
            # to reduce the computation we are
            # dictionary[word] = idf value of word in whole corpus
            # sent.count(word) = tf value of word in this review
            tf_idf = dictionary[word]*(sent.count(word)/len(sent))
            sent_vec += (vec * tf_idf)
            weight_sum += tf_idf
    if weight_sum != 0:
        sent_vec /= weight_sum
    tfidf_sent_vectorsPT_train.append(sent_vec)
    row += 1

100%| 33500/33500 [00:22<00:00, 1503.39it/s]

```

Test

```

In [0]: # S = ["abc def pqr", "def def def abc", "pqr pqr def"]
model = TfidfVectorizer()
model.fit(X_test['clean_project_title'])
# we are converting a dictionary with word as a key, and the idf as a value
dictionary = dict(zip(model.get_feature_names(), list(model.idf_)))

In [61]: # TF-IDF weighted Word2Vec
tfidf_feat = model.get_feature_names() # tfidf words/col-names
# final_tf_idf is the sparse matrix with row= sentence, col=word and cell_val = tfidf

tfidf_sent_vectorsPT_test = []; # the tfidf-w2v for each sentence/review is stored in this list
row=0;
for sent in tqdm(list_of_sentencePT_test): # for each review/sentence
    sent_vec = np.zeros(50) # as word vectors are of zero length
    weight_sum =0; # num of words with a valid vector in the sentence/review
    for word in sent: # for each word in a review/sentence
        if word in w2v_words and word in tfidf_feat:
            vec = w2v_model.wv[word]
#            tf_idf = tf_idf_matrix[row, tfidf_feat.index(word)]

```

```

# to reduce the computation we are
# dictionary[word] = idf value of word in whole corpus
# sent.count(word) = tf value of word in this review
tf_idf = dictionary[word]*(sent.count(word)/len(sent))
sent_vec += (vec * tf_idf)
weight_sum += tf_idf
if weight_sum != 0:
    sent_vec /= weight_sum
tfidf_sent_vectorsPT_test.append(sent_vec)
row += 1

```

100%| 16500/16500 [00:08<00:00, 1905.55it/s]

1.6.15 1.5.3 Vectorizing Numerical features

1.6.16 price

```

In [0]: price_data = resource_data.groupby('id').agg({'price':'sum', 'quantity':'sum'}).reset_index()
X_train = pd.merge(X_train, price_data, on='id', how='left')
#X_cv = pd.merge(X_cv, price_data, on='id', how='left')
X_test = pd.merge(X_test, price_data, on='id', how='left')

```

In [63]: X_train.head(1)

```

Out[63]:  Unnamed: 0      id      teacher_id teacher_prefix \
0      36401  p111136  a703366918b87aaa7eaf013321ccd98f      Mrs.

      school_state      Date project_grade_category \
0      GA 2016-08-22 12:53:43      Grades 6-8

      project_essay_1 \
0  We are a Montessori school with a deep history...

      project_essay_2 project_essay_3 \
0  We are moving to a healthier student body and ...      NaN

      project_essay_4      project_resource_summary \
0      NaN  My students need some fitness stuff.  We are m...

      teacher_number_of_previously_posted_projects clean_categories \
0      5      Health_Sports

      clean_subcategories      clean_essay \
0      Gym_Fitness  montessori school deep history student centere...

      clean_project_title  price  quantity
0      fitness frenzy  613.96      15

```

```
In [64]: from sklearn.preprocessing import Normalizer
normalizer = Normalizer()
# normalizer.fit(X_train['price'].values)
# this will rise an error Expected 2D array, got 1D array instead:
# array=[105.22 215.96 96.01 ... 368.98 80.53 709.67].
# Reshape your data either using
# array.reshape(-1, 1) if your data has a single feature
# array.reshape(1, -1) if it contains a single sample.
normalizer.fit(X_train['price'].values.reshape(-1,1))

X_train_price_norm = normalizer.transform(X_train['price'].values.reshape(-1,1))
#X_cv_price_norm = normalizer.transform(X_cv['price'].values.reshape(-1,1))
X_test_price_norm = normalizer.transform(X_test['price'].values.reshape(-1,1))

print("After vectorizations")
print(X_train_price_norm.shape, y_train.shape)
#print(X_cv_price_norm.shape, y_cv.shape)
print(X_test_price_norm.shape, y_test.shape)
print("="*100)
```

After vectorizations

(33500, 1) (33500,)

(16500, 1) (16500,)

=====

1.6.17 tnppp

```
In [65]: from sklearn.preprocessing import Normalizer
normalizerT = Normalizer()
# normalizer.fit(X_train['price'].values)
# this will rise an error Expected 2D array, got 1D array instead:
# array=[105.22 215.96 96.01 ... 368.98 80.53 709.67].
# Reshape your data either using
# array.reshape(-1, 1) if your data has a single feature
# array.reshape(1, -1) if it contains a single sample.
normalizerT.fit(X_train['teacher_number_of_previously_posted_projects'].values.reshape(-1,1))

X_train_tnppp_norm = normalizerT.transform(X_train['teacher_number_of_previously_posted_projects'].values.reshape(-1,1))
#X_cv_tnppp_norm = normalizerT.transform(X_cv['teacher_number_of_previously_posted_projects'].values.reshape(-1,1))
X_test_tnppp_norm = normalizerT.transform(X_test['teacher_number_of_previously_posted_projects'].values.reshape(-1,1))

print("After vectorizations")
print(X_train_tnppp_norm.shape, y_train.shape)
#print(X_cv_tnppp_norm.shape, y_cv.shape)
print(X_test_tnppp_norm.shape, y_test.shape)
print("="*100)
```

After vectorizations

```
(33500, 1) (33500,)  
(16500, 1) (16500,)  
=====
```

1.6.18 1.5.4 Merging all the above features

- we need to merge all the numerical vectors i.e catogorical, text, numerical vectors

SET1

In [66]: # merge two sparse matrices: <https://stackoverflow.com/a/19710648/4084039>

```
from scipy.sparse import hstack  
X_tr1 = hstack((X_train_cc_ohe, X_train_csc_ohe, X_train_grade_ohe, X_train_price_norm, X_train_tnp  
#X_cr1 = hstack((X_cv_cc_ohe, X_cv_csc_ohe, X_cv_grade_ohe, X_cv_price_norm, X_cv_tnp  
X_te1 = hstack((X_test_cc_ohe, X_test_csc_ohe, X_test_grade_ohe, X_test_price_norm, X_test_t  
  
print("Final Data matrix")  
print(X_tr1.shape, y_train.shape)  
#print(X_cr1.shape, y_cv.shape)  
print(X_te1.shape, y_test.shape)  
print("="*100)
```

Final Data matrix

```
(33500, 6713) (33500,)  
(16500, 6713) (16500,)  
=====
```

SET2

In [67]: # merge two sparse matrices: <https://stackoverflow.com/a/19710648/4084039>

```
from scipy.sparse import hstack  
X_tr2 = hstack((X_train_cc_ohe, X_train_csc_ohe, X_train_grade_ohe, X_train_price_norm, X_train_tnp  
#X_cr2 = hstack((X_cv_cc_ohe, X_cv_csc_ohe, X_cv_grade_ohe, X_cv_price_norm, X_cv_tnp  
X_te2 = hstack((X_test_cc_ohe, X_test_csc_ohe, X_test_grade_ohe, X_test_price_norm, X_test_t  
  
print("Final Data matrix")  
print(X_tr2.shape, y_train.shape)  
#print(X_cr2.shape, y_cv.shape)  
print(X_te2.shape, y_test.shape)  
print("="*100)
```

Final Data matrix

```
(33500, 6713) (33500,)  
(16500, 6713) (16500,)  
=====
```

SET3

In [68]: # merge two sparse matrices: <https://stackoverflow.com/a/19710648/4084039>

```
from scipy.sparse import hstack
X_tr3 = hstack((X_train_cc_ohe, X_train_csc_ohe, X_train_grade_ohe, X_train_price_norm, X_train_tnpnp_ohe))
#X_cr3 = hstack((X_cv_cc_ohe, X_cv_csc_ohe, X_cv_grade_ohe, X_cv_price_norm, X_cv_tnpnp_ohe))
X_te3 = hstack((X_test_cc_ohe, X_test_csc_ohe, X_test_grade_ohe, X_test_price_norm, X_test_tnpnp_ohe))

print("Final Data matrix")
print(X_tr3.shape, y_train.shape)
#print(X_cr3.shape, y_cv.shape)
print(X_te3.shape, y_test.shape)
print("="*100)
```

Final Data matrix

(33500, 144) (33500,)

(16500, 144) (16500,)

=====

SET4

In [69]: # merge two sparse matrices: <https://stackoverflow.com/a/19710648/4084039>

```
from scipy.sparse import hstack
X_tr4 = hstack((X_train_cc_ohe, X_train_csc_ohe, X_train_grade_ohe, X_train_price_norm, X_train_tnpnp_ohe))
#X_cr4 = hstack((X_cv_cc_ohe, X_cv_csc_ohe, X_cv_grade_ohe, X_cv_price_norm, X_cv_tnpnp_ohe))
X_te4 = hstack((X_test_cc_ohe, X_test_csc_ohe, X_test_grade_ohe, X_test_price_norm, X_test_tnpnp_ohe))

print("Final Data matrix")
print(X_tr4.shape, y_train.shape)
#print(X_cr4.shape, y_cv.shape)
print(X_te4.shape, y_test.shape)
print("="*100)
```

Final Data matrix

(33500, 144) (33500,)

(16500, 144) (16500,)

=====

1.6.19 Computing Sentiment Scores

```
In [0]: essay_List_tr = X_train['clean_essay']
        essay_List_te = X_test['clean_essay']
```

Train

```
In [71]: import nltk
         from nltk.sentiment.vader import SentimentIntensityAnalyzer
```

```
nltk.download('vader_lexicon')
```

```
sia = SentimentIntensityAnalyzer()
```

```
ss_tr = []
ss_neg_tr = []
ss_neu_tr = []
ss_pos_tr = []
ss_com_tr = []

for i in range(len(essay_List_tr)) :
    ss_tr.append(sia.polarity_scores(essay_List_tr[i]))
    ss_neg_tr.append(ss_tr[i]['neg'])
    ss_neu_tr.append(ss_tr[i]['neu'])
    ss_pos_tr.append(ss_tr[i]['pos'])
    ss_com_tr.append(ss_tr[i]['compound'])
```

```
[nltk_data] Downloading package vader_lexicon to /root/nltk_data...
```

```
[nltk_data] Package vader_lexicon is already up-to-date!
```

```
In [0]: X_train['negative_sentiment_score'] = ss_neg_tr
        X_train['neutral_sentiment_score'] = ss_neu_tr
        X_train['positive_sentiment_score'] = ss_pos_tr
        X_train['compound_sentiment_score'] = ss_com_tr
```

Test

```
In [0]: import nltk
        from nltk.sentiment.vader import SentimentIntensityAnalyzer
```

```
sia = SentimentIntensityAnalyzer()
```

```
ss_te = []
ss_neg_te = []
ss_neu_te = []
ss_pos_te = []
ss_com_te = []

for i in range(len(essay_List_te)) :
    ss_te.append(sia.polarity_scores(essay_List_te[i]))
    ss_neg_te.append(ss_te[i]['neg'])
    ss_neu_te.append(ss_te[i]['neu'])
    ss_pos_te.append(ss_te[i]['pos'])
    ss_com_te.append(ss_te[i]['compound'])
```

```
In [0]: X_test['negative_sentiment_score'] = ss_neg_te
        X_test['neutral_sentiment_score'] = ss_neu_te
```



```
X_test['positive_sentiment_score'] = ss_pos_te
X_test['compound_sentiment_score'] = ss_com_te
```

1.6.20 Computing no of words in essay

```
In [0]: def num_Words(s) :
        return len(s.split())
```

Train

```
In [0]: num_wo_tr = []
        for i in X_train['clean_essay'] :
            num_wo_tr.append(num_Words(i))
```

```
In [0]: X_train['num_words_essay'] = num_wo_tr
```

Test

```
In [0]: num_wo_te = []
        for i in X_test['clean_essay'] :
            num_wo_te.append(num_Words(i))
```

```
In [0]: X_test['num_words_essay'] = num_wo_te
```

1.6.21 Computing number of words in project_title

Train

```
In [0]: num_wopt_tr = []
        for i in X_train['clean_project_title'] :
            num_wopt_tr.append(num_Words(i))
```

```
In [0]: X_train['num_words_project_title'] = num_wopt_tr
```

Test

```
In [0]: num_wopt_te = []
        for i in X_test['clean_project_title'] :
            num_wopt_te.append(num_Words(i))
```

```
In [0]: X_test['num_words_project_title'] = num_wopt_te
```

```
In [84]: X_test.head(1)
```

```
Out[84]:  Unnamed: 0      id      teacher_id teacher_prefix \
0      91315  p241375  c47f1bc58eeee0ce9831f77f14dcc892      Mrs.

      school_state      Date project_grade_category \
0      TX 2016-09-06 17:00:13      Grades PreK-2
```

```

                                project_essay_1 \
0  My class is always moving. We work on projects...

                                project_essay_2 project_essay_3 \
0  I envision a room of student-led choice, where...      NaN

                                clean_essay \
0  ... class always moving work projects solve proble...

                                clean_project_title  price quantity negative_sentiment_score \
0  ipads student led classroom  267.74      12      0.074

                                neutral_sentiment_score positive_sentiment_score compound_sentiment_score \
0                                0.622      0.304      0.9862

                                num_words_essay num_words_project_title
0                                136      4

[1 rows x 25 columns]

```

2 Assignment 5: Logistic Regression

[Task-1] Logistic Regression(either SGDClassifier with log loss, or LogisticRegression) on these features

- Set 1: categorical, numerical features + project_title(BOW) + preprocessed_grams` with `min_df=10` and `max_features=5000`)
- Set 2: categorical, numerical features + project_title(TFIDF)+ preprocessed_grams` with `min_df=10` and `max_features=5000`)
- Set 3: categorical, numerical features + project_title(AVG W2V)+ preprocessed_grams` with `min_df=10` and `max_features=5000`)
- Set 4: categorical, numerical features + project_title(TFIDF W2V)+ preprocessed_grams` with `min_df=10` and `max_features=5000`)

Hyper paramter tuning (find best hyper parameters corresponding the algorithm that you choose)

- Find the best hyper parameter which will give the maximum AUC value
- Find the best hyper paramter using k-fold cross validation or simple cross validation data
- Use gridsearch cv or randomsearch cv or you can also write your own for loops to do this task of hyperparameter

Representation of results

- You need to plot the performance of model both on train data and cross validation data for each hyper parameter

Once after you found the best hyper parameter, you need to train your model with it, and find the AUC on test data

Along with plotting ROC curve, you need to print the confusion matrix with predicted and original labels

[Task-2] Apply Logistic Regression on the below feature set Set 5 by finding the best hyper parameter

 Consider these set of features Set 5 :

school_state : categorical data

clean_categories : categorical data

clean_subcategories : categorical data

project_grade_category :categorical data

teacher_prefix : categorical data

quantity : numerical data

teacher_number_of_previously_posted_projects : numerical data

price : numerical data

sentiment score's of each of the essay : numerical data

number of words in the title : numerical data

number of words in the combine essays : numerical data

And apply the Logistic regression on these features by finding the best hyper parameter as suggested in step 2 and step 3

Conclusion

You need to summarize the results at the end of the notebook, summarize it in the table format. To print out a table

Note: Data Leakage

1. There will be an issue of data-leakage if you vectorize the entire data and then split it into train/cv/test.
2. To avoid the issue of data-leakage, make sure to split your data first and then vectorize it.
3. While vectorizing your data, apply the method fit_transform() on your train data, and apply the method transform() on cv/test data.
4. For more details please go through this link.

2. Logistic Regression

2.4 Applying Logistic Regression on different kind of featurization as mentioned in the instructions

Apply Logistic Regression on different kind of featurization as mentioned in the instructions
For Every model that you work on make sure you do the step 2 and step 3 of instructions

2.0.1 SET1

```
In [103]: # https://scikit-learn.org/stable/modules/generated/sklearn.model\_selection.GridSearchCV.html
from sklearn.model_selection import GridSearchCV
from sklearn.linear_model import LogisticRegression

lr1 = LogisticRegression()

parameters = {'C': [10**-4, 10**-3, 10**-2, 10**-1, 10**0, 10**1, 10**2, 10**3, 10**4]}
#parameters = {'alpha':[0.0001, 0.0025, 0.0005, 0.0075, 0.001, 0.025, .005, 0.075, 0.1, 0.25, 0.5, 0.75, 1]}

clf = GridSearchCV(lr1, parameters, cv=3, scoring='roc_auc')
clr1 = clf.fit(X_tr1, y_train)

train_auc1= clf.cv_results_['mean_train_score']
train_auc_std1= clf.cv_results_['std_train_score']
cv_auc1 = clf.cv_results_['mean_test_score']
cv_auc_std1 = clf.cv_results_['std_test_score']

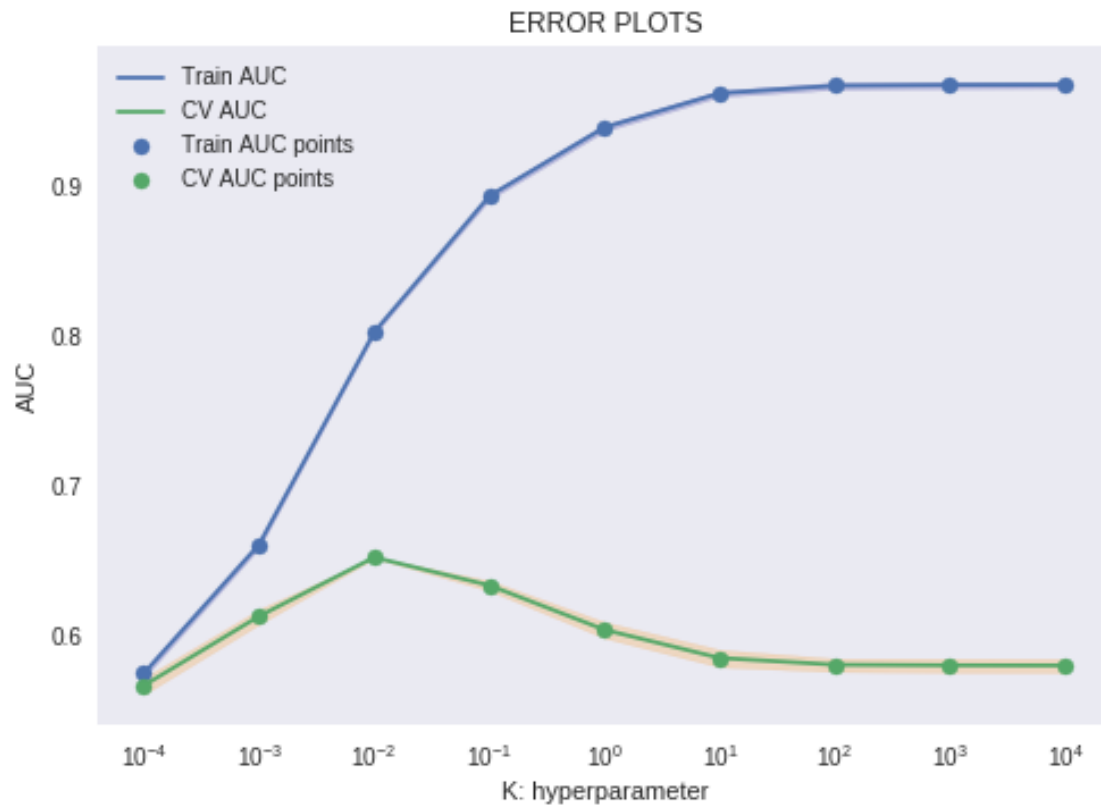
plt.plot(parameters['C'], train_auc1, label='Train AUC')
# this code is copied from here: https://stackoverflow.com/a/48803361/4084039
plt.gca().fill_between(parameters['C'],train_auc1 - train_auc_std1,train_auc1 + train_auc_std1,alpha=0.2,color='red')

plt.plot(parameters['C'], cv_auc1, label='CV AUC')
# this code is copied from here: https://stackoverflow.com/a/48803361/4084039
plt.gca().fill_between(parameters['C'],cv_auc1 - cv_auc_std1, cv_auc1 + cv_auc_std1,alpha=0.2,color='red')

plt.scatter(parameters['C'], train_auc1, label='Train AUC points')
plt.scatter(parameters['C'], cv_auc1, label='CV AUC points')

plt.xscale('log')
plt.legend()
plt.xlabel("K: hyperparameter")
plt.ylabel("AUC")
plt.title("ERROR PLOTS")
plt.grid()
plt.show()

print(train_auc1)
print(cv_auc1)      #C=10**-2
```



```
[0.57481995 0.66096287 0.80286834 0.8939249 0.93984093 0.96267749
 0.96773145 0.96818678 0.96834748]
[0.56667613 0.61297126 0.65232766 0.63366294 0.6039576 0.58509715
 0.58059071 0.58020222 0.58006535]
```

2.0.2 SET2

```
In [86]: # https://scikit-learn.org/stable/modules/generated/sklearn.model_selection.GridSearchCV.html
from sklearn.model_selection import GridSearchCV
from sklearn.linear_model import LogisticRegression

lr1 = LogisticRegression()

parameters = {'C': [10**-4, 10**-3, 10**-2, 10**-1, 10**0, 10**1, 10**2, 10**3, 10**4]}
#parameters = {'alpha':[0.0001, 0.0025, 0.0005, 0.0075, 0.001, 0.025, .005, 0.075, 0.1, 0.25, 0.5, 0.75, 1]}

clf = GridSearchCV(lr1, parameters, cv=5, scoring='roc_auc')
clr2 = clf.fit(X_tr2, y_train)

train_auc2= clf.cv_results_['mean_train_score']
```

```

train_auc_std2= clf.cv_results_['std_train_score']
cv_auc2 = clf.cv_results_['mean_test_score']
cv_auc_std2 = clf.cv_results_['std_test_score']

plt.plot(parameters['C'], train_auc2, label='Train AUC')
# this code is copied from here: https://stackoverflow.com/a/48803361/4084039
plt.gca().fill_between(parameters['C'],train_auc2 - train_auc_std2,train_auc2 + train_auc_std2,alpha=0.2,color='blue')


plt.plot(parameters['C'], cv_auc2, label='CV AUC')
# this code is copied from here: https://stackoverflow.com/a/48803361/4084039
plt.gca().fill_between(parameters['C'],cv_auc2 - cv_auc_std2, cv_auc2 + cv_auc_std2,alpha=0.2,color='green')

plt.scatter(parameters['C'], train_auc2, label='Train AUC points')
plt.scatter(parameters['C'], cv_auc2, label='CV AUC points')

plt.xscale('log')
plt.legend()
plt.xlabel("K: hyperparameter")
plt.ylabel("AUC")
plt.title("ERROR PLOTS")
plt.grid()
plt.show()

print(train_auc2)
print(cv_auc2)      #k_best = 1.0

```



```
[0.53784526 0.55847431 0.6199883  0.75369815 0.86473428 0.9142799
 0.92911808 0.93209242 0.93254943]
[0.5370075  0.55298354 0.5904785  0.64337824 0.64321038 0.61262591
 0.59236436 0.58517655 0.58301739]
```

2.0.3 SET3

```
In [87]: # https://scikit-learn.org/stable/modules/generated/sklearn.model\_selection.GridSearchCV.html
from sklearn.model_selection import GridSearchCV
from sklearn.linear_model import LogisticRegression

lr1 = LogisticRegression()

parameters = {'C': [10**-4, 10**-3, 10**-2, 10**-1, 10**0, 10**1, 10**2, 10**3, 10**4]}
#parameters = {'alpha':[0.0001, 0.0025, 0.0005, 0.0075, 0.001, 0.025, .005, 0.075, 0.1, 0.25, 0.5, 0.75, 1]}

clf = GridSearchCV(lr1, parameters, cv=5, scoring='roc_auc')
clr3 = clf.fit(X_tr3, y_train)

train_auc3= clf.cv_results_['mean_train_score']
train_auc_std3= clf.cv_results_['std_train_score']
cv_auc3 = clf.cv_results_['mean_test_score']
cv_auc_std3 = clf.cv_results_['std_test_score']

plt.plot(parameters['C'], train_auc3, label='Train AUC')
# this code is copied from here: https://stackoverflow.com/a/48803361/4084039
plt.gca().fill_between(parameters['C'],train_auc3 - train_auc_std3,train_auc3 + train_auc_std3,alpha=0.2,color='red')

plt.plot(parameters['C'], cv_auc3, label='CV AUC')
# this code is copied from here: https://stackoverflow.com/a/48803361/4084039
plt.gca().fill_between(parameters['C'],cv_auc3 - cv_auc_std3, cv_auc3 + cv_auc_std3,alpha=0.2,color='red')

plt.scatter(parameters['C'], train_auc3, label='Train AUC points')
plt.scatter(parameters['C'], cv_auc3, label='CV AUC points')

plt.xscale('log')
plt.legend()
plt.xlabel("K: hyperparameter")
plt.ylabel("AUC")
plt.title("ERROR PLOTS")
plt.grid()
plt.show()

print(train_auc3)
print(cv_auc3)      #k_best = 1.0
```



```
[0.51357013 0.61940546 0.67103939 0.68094688 0.68354276 0.68530282
 0.68580518 0.68583099 0.68583481]
[0.51290845 0.61587312 0.66380244 0.67003152 0.67045306 0.67090011
 0.6706912  0.6706031  0.67060187]
```

2.0.4 SET4

```
In [88]: # https://scikit-learn.org/stable/modules/generated/sklearn.model_selection.GridSearchCV.html
from sklearn.model_selection import GridSearchCV
from sklearn.linear_model import LogisticRegression

lr1 = LogisticRegression()

parameters = {'C': [10**-4, 10**-3, 10**-2, 10**-1, 10**0, 10**1, 10**2, 10**3, 10**4]}
#parameters = {'alpha':[0.0001, 0.0025, 0.0005, 0.0075, 0.001, 0.025, .005, 0.075, 0.1, 0.25, 0.5, 0.75, 1]}

clf = GridSearchCV(lr1, parameters, cv=5, scoring='roc_auc')
clr4 = clf.fit(X_tr4, y_train)

train_auc4= clf.cv_results_['mean_train_score']
train_auc_std4= clf.cv_results_['std_train_score']
cv_auc4 = clf.cv_results_['mean_test_score']
```



```

cv_auc_std4 = clf.cv_results_['std_test_score']

plt.plot(parameters['C'], train_auc4, label='Train AUC')
# this code is copied from here: https://stackoverflow.com/a/48803361/4084039
plt.gca().fill_between(parameters['C'],train_auc4 - train_auc_std4,train_auc4 + train_auc_std4,alpha=0.2,color=


plt.plot(parameters['C'], cv_auc4, label='CV AUC')
# this code is copied from here: https://stackoverflow.com/a/48803361/4084039
plt.gca().fill_between(parameters['C'],cv_auc4 - cv_auc_std4, cv_auc4 + cv_auc_std4,alpha=0.2,color=

plt.scatter(parameters['C'], train_auc4, label='Train AUC points')
plt.scatter(parameters['C'], cv_auc4, label='CV AUC points')

plt.xscale('log')
plt.legend()
plt.xlabel("K: hyperparameter")
plt.ylabel("AUC")
plt.title("ERROR PLOTS")
plt.grid()
plt.show()

print(train_auc4)
print(cv_auc4)      #k_best = 10**4

```



```
[0.51706877 0.53906059 0.56580373 0.57472062 0.58282907 0.60414265
 0.6237728  0.62893627 0.62931895]
[0.51678628 0.53711483 0.55969416 0.56481471 0.569221  0.5844401
 0.60041664 0.60528599 0.60558444]
```

2.0.5 Plot of AUC on Test and Train

SET1

```
In [0]: def batch_predict(clf, data):
        # roc_auc_score(y_true, y_score) the 2nd parameter should be probability estimates of the positive class
        # not the predicted outputs

        y_data_pred = []
        tr_loop = data.shape[0] - data.shape[0]%1000
        # consider you X_tr shape is 49041, then your cr_loop will be 49041 - 49041%1000 = 49000
        # in this for loop we will iterate until the last 1000 multiplier
        for i in range(0, tr_loop, 1000):
            y_data_pred.extend(clf.predict_proba(data[i:i+1000])[:,1])
        # we will be predicting for the last data points
        y_data_pred.extend(clf.predict_proba(data[tr_loop:])[:,1])

        return y_data_pred

In [105]: # https://scikit-learn.org/stable/modules/generated/sklearn.metrics.roc_curve.html#sklearn.metrics.roc_curve

from sklearn.linear_model import LogisticRegression
from sklearn.metrics import roc_curve, auc


clf = LogisticRegression(C=10**-2, penalty='l1');
clf.fit(X_tr1, y_train)
# roc_auc_score(y_true, y_score) the 2nd parameter should be probability estimates of the positive class
# not the predicted outputs

y_train_pred1 = batch_predict(clf, X_tr1)
y_test_pred1 = batch_predict(clf, X_te1)

train_fpr1, train_tpr1, tr_thresholds1 = roc_curve(y_train, y_train_pred1)
test_fpr1, test_tpr1, te_thresholds1 = roc_curve(y_test, y_test_pred1)

plt.plot(train_fpr1, train_tpr1, label="train AUC =" + str(auc(train_fpr1, train_tpr1)))
plt.plot(test_fpr1, test_tpr1, label="test AUC =" + str(auc(test_fpr1, test_tpr1)))
plt.legend()
plt.xlabel("False Positive Rate")
plt.ylabel("True Positive Rate")
plt.title("ERROR PLOTS")
```

```
plt.grid()
plt.show()
```



Confusion Matrix

```
In [0]: def predict(proba, threshold, fpr, tpr):
```

```
    t = threshold[np.argmax(fpr*(1-tpr))]
```

```
    # (tpr*(1-fpr)) will be maximum if your fpr is very low and tpr is very high
```

```
    print("the maximum value of tpr*(1-fpr)", max(tpr*(1-fpr)), "for threshold", np.round(t,3))
```

```
    predictions = []
```

```
    for i in proba:
```

```
        if i >= t:
```

```
            predictions.append(1)
```

```
        else:
```

```
            predictions.append(0)
```

```
    return predictions
```

```
In [107]: #https://www.quantinsti.com/blog/creating-heatmap-using-python-seaborn
```

```
import seaborn as sns; sns.set()
```

```

con_m_train = confusion_matrix(y_train, predict(y_train_pred1, tr_thresholds1, train_fpr1, train_tpr1))
con_m_test = confusion_matrix(y_test, predict(y_test_pred1, te_thresholds1, test_fpr1, test_tpr1))

key = (np.asarray([['TN', 'FP'], ['FN', 'TP']]))

fig, ax = plt.subplots(1,2, figsize=(12,5))

labels_train = (np.asarray(["{0} = {1:.2f}" .format(key, value) for key, value in zip(key.flatten(), con_m_train.flatten())])
labels_test = (np.asarray(["{0} = {1:.2f}" .format(key, value) for key, value in zip(key.flatten(), con_m_test.flatten())])


sns.heatmap(con_m_train, linewidths=.5, xticklabels=['PREDICTED : NO', 'PREDICTED : YES'], yticklabels=['ACTUAL : NO', 'ACTUAL : YES'])
sns.heatmap(con_m_test, linewidths=.5, xticklabels=['PREDICTED : NO', 'PREDICTED : YES'], yticklabels=['ACTUAL : NO', 'ACTUAL : YES'])

ax[0].set_title('Train Set')
ax[1].set_title('Test Set')

plt.show()

```

the maximum value of $tpr*(1-fpr)$ 0.3109537334628896 for threshold 0.835
the maximum value of $tpr*(1-fpr)$ 0.32179050659213454 for threshold 0.833



SET2

In [134]: # https://scikit-learn.org/stable/modules/generated/sklearn.metrics.roc_curve.html#sklearn.metrics.roc_curve

```
from sklearn.metrics import roc_curve, auc
```

```


clf = LogisticRegression(C=1, penalty='l1');
clf.fit(X_tr2, y_train)
# roc_auc_score(y_true, y_score) the 2nd parameter should be probability estimates of the positive class
# not the predicted outputs

y_train_pred2 = batch_predict(clf, X_tr2)
y_test_pred2 = batch_predict(clf, X_te2)

train_fpr2, train_tpr2, tr_thresholds2 = roc_curve(y_train, y_train_pred2)
test_fpr2, test_tpr2, te_thresholds2 = roc_curve(y_test, y_test_pred2)

plt.plot(train_fpr2, train_tpr2, label="train AUC =" + str(auc(train_fpr2, train_tpr2)))
plt.plot(test_fpr2, test_tpr2, label="test AUC =" + str(auc(test_fpr2, test_tpr2)))
plt.legend()
plt.xlabel("False Positive Rate")
plt.ylabel("True Positive Rate")
plt.title("ERROR PLOTS")
plt.grid()
plt.show()

```



Confusion Matrix

In [0]: `def predict(proba, threshold, fpr, tpr):`

```
t = threshold[np.argmax(fpr*(1-tpr))]

# (tpr*(1-fpr)) will be maximum if your fpr is very low and tpr is very high

print("the maximum value of tpr*(1-fpr)", max(tpr*(1-fpr)), "for threshold", np.round(t,3))
predictions = []
for i in proba:
    if i>=t:
        predictions.append(1)
    else:
        predictions.append(0)
return predictions
```

In [136]: `#https://www.quantinsti.com/blog/creating-heatmap-using-python-seaborn`

```
import seaborn as sns; sns.set()
```

```
con_m_train = confusion_matrix(y_train, predict(y_train_pred2, tr_thresholds2, train_fpr2, train_tpr2))
con_m_test = confusion_matrix(y_test, predict(y_test_pred2, te_thresholds2, test_fpr2, test_tpr2))
```

```
key = (np.asarray([['TN', 'FP'], ['FN', 'TP']]))
```

```
fig, ax = plt.subplots(1,2, figsize=(12,5))
```

```
labels_train = (np.asarray(["{0} = {1:.2f}" .format(key, value) for key, value in zip(key.flatten(), con_m_train.flatten())])
labels_test = (np.asarray(["{0} = {1:.2f}" .format(key, value) for key, value in zip(key.flatten(), con_m_test.flatten())])
```

```
sns.heatmap(con_m_train, linewidths=.5, xticklabels=['PREDICTED : NO', 'PREDICTED : YES'], yticklabels=labels_train)
sns.heatmap(con_m_test, linewidths=.5, xticklabels=['PREDICTED : NO', 'PREDICTED : YES'], yticklabels=labels_test)
```


```
ax[0].set_title('Train Set')
```

```
ax[1].set_title('Test Set')
```

```
plt.show()
```

the maximum value of $tpr*(1-fpr)$ 0.533311115023151 for threshold 0.836

the maximum value of $tpr*(1-fpr)$ 0.3794445355804354 for threshold 0.85



SET3

In [111]: # https://scikit-learn.org/stable/modules/generated/sklearn.metrics.roc_curve.html#sklearn.metrics.roc_curve


```
from sklearn.metrics import roc_curve, auc

clf = LogisticRegression(C=1, penalty='l1');
clf.fit(X_tr3, y_train)
# roc_auc_score(y_true, y_score) the 2nd parameter should be probability estimates of the positive class
# not the predicted outputs

y_train_pred3 = batch_predict(clf, X_tr3)
y_test_pred3 = batch_predict(clf, X_te3)

train_fpr3, train_tpr3, tr_thresholds3 = roc_curve(y_train, y_train_pred3)
test_fpr3, test_tpr3, te_thresholds3 = roc_curve(y_test, y_test_pred3)

plt.plot(train_fpr3, train_tpr3, label="train AUC =" + str(auc(train_fpr3, train_tpr3)))
plt.plot(test_fpr3, test_tpr3, label="test AUC =" + str(auc(test_fpr3, test_tpr3)))
plt.legend()
plt.xlabel("False Positive Rate")
plt.ylabel("True Positive Rate")
plt.title("ERROR PLOTS")
plt.grid()
plt.show()
```



Confusion Matrix

In [0]: `def predict(proba, threshold, fpr, tpr):`

```
    t = threshold[np.argmax(fpr*(1-tpr))]
```

```
    # (tpr*(1-fpr)) will be maximum if your fpr is very low and tpr is very high
```

```
    print("the maximum value of tpr*(1-fpr)", max(tpr*(1-fpr)), "for threshold", np.round(t,3))
```

```
    predictions = []
```

```
    for i in proba:
```

```
        if i >= t:
```

```
            predictions.append(1)
```

```
        else:
```

```
            predictions.append(0)
```

```
    return predictions
```

In [113]: `#https://www.quantinsti.com/blog/creating-heatmap-using-python-seaborn`

```
import seaborn as sns; sns.set()
```

```
con_m_train = confusion_matrix(y_train, predict(y_train_pred3, tr_thresholds3, train_fpr3, train_tpr3))
```

```
con_m_test = confusion_matrix(y_test, predict(y_test_pred3, te_thresholds3, test_fpr3, test_tpr3))
```



```

key = (np.asarray([['TN', 'FP'], ['FN', 'TP']]))

fig, ax = plt.subplots(1, 2, figsize=(12, 5))

labels_train = (np.asarray(["{0} = {1:.2f}" .format(key, value) for key, value in zip(key.flatten(), con_m_train)])
labels_test = (np.asarray(["{0} = {1:.2f}" .format(key, value) for key, value in zip(key.flatten(), con_m_test)])


sns.heatmap(con_m_train, linewidths=.5, xticklabels=['PREDICTED : NO', 'PREDICTED : YES'], yticklabels=labels_train)
sns.heatmap(con_m_test, linewidths=.5, xticklabels=['PREDICTED : NO', 'PREDICTED : YES'], yticklabels=labels_test)

ax[0].set_title('Train Set')
ax[1].set_title('Test Set')

plt.show()

```

the maximum value of $tpr * (1 - fpr)$ 0.40126268189346226 for threshold 0.846
the maximum value of $tpr * (1 - fpr)$ 0.32493893919883315 for threshold 0.994



SET4

In [114]: # https://scikit-learn.org/stable/modules/generated/sklearn.metrics.roc_curve.html#sklearn.metrics.roc_curve

```

from sklearn.metrics import roc_curve, auc

clf = LogisticRegression(C=10**4, penalty='l1');
clf.fit(X_tr4, y_train)


```

```
# roc_auc_score(y_true, y_score) the 2nd parameter should be probability estimates of the positive class
# not the predicted outputs
```

```
y_train_pred4 = batch_predict(clf, X_tr4)
y_test_pred4 = batch_predict(clf, X_te4)
```

```
train_fpr4, train_tpr4, tr_thresholds4 = roc_curve(y_train, y_train_pred4)
test_fpr4, test_tpr4, te_thresholds4 = roc_curve(y_test, y_test_pred4)
```

```
plt.plot(train_fpr4, train_tpr4, label="train AUC =" + str(auc(train_fpr4, train_tpr4)))
plt.plot(test_fpr4, test_tpr4, label="test AUC =" + str(auc(test_fpr4, test_tpr4)))
plt.legend()
plt.xlabel("False Positive Rate")
plt.ylabel("True Positive Rate")
plt.title("ERROR PLOTS")
plt.grid()
plt.show()
```



Confusion Matrix

```
In [0]: def predict(proba, threshold, fpr, tpr):
```

```

t = threshold[np.argmax(fpr*(1-tpr))]

# (tpr*(1-fpr)) will be maximum if your fpr is very low and tpr is very high

print("the maximum value of tpr*(1-fpr)", max(tpr*(1-fpr)), "for threshold", np.round(t,3))
predictions = []
for i in proba:
    if i>=t:
        predictions.append(1)
    else:
        predictions.append(0)
return predictions

```

In [116]: [#https://www.quantinsti.com/blog/creating-heatmap-using-python-seaborn](https://www.quantinsti.com/blog/creating-heatmap-using-python-seaborn)
import seaborn as sns; sns.set()

```

con_m_train = confusion_matrix(y_train, predict(y_train_pred4, tr_thresholds4, train_fpr4, train_tpr4))
con_m_test = confusion_matrix(y_test, predict(y_test_pred4, te_thresholds4, test_fpr4, test_tpr4))

key = (np.asarray([['TN','FP'], ['FN', 'TP']]))

fig, ax = plt.subplots(1,2, figsize=(12,5))

labels_train = (np.asarray(["{0} = {1:.2f}" .format(key, value) for key, value in zip(key.flatten(), con_m_train.flatten())])
labels_test = (np.asarray(["{0} = {1:.2f}" .format(key, value) for key, value in zip(key.flatten(), con_m_test.flatten())])

sns.heatmap(con_m_train, linewidths=.5, xticklabels=['PREDICTED : NO', 'PREDICTED : YES'], yticklabels=labels_train)
sns.heatmap(con_m_test, linewidths=.5, xticklabels=['PREDICTED : NO', 'PREDICTED : YES'], yticklabels=labels_test)


ax[0].set_title('Train Set')
ax[1].set_title('Test Set')

plt.show()

```

the maximum value of tpr*(1-fpr) 0.34427297580348704 for threshold 0.845

the maximum value of tpr*(1-fpr) 0.3372308137127929 for threshold 0.846



2.5 Logistic Regression with added Features Set 5

2.1 Encoding additional numeric features

In [102]: X_train.head(1)

```
Out[102]: Unnamed: 0      id      teacher_id teacher_prefix \
0      36401  p111136  a703366918b87aaa7eaf013321ccd98f      Mrs.

      school_state      Date project_grade_category \
0      GA 2016-08-22 12:53:43      Grades 6-8

      project_essay_1 \
0  We are a Montessori school with a deep history...

      project_essay_2 project_essay_3 \
0  We are moving to a healthier student body and ...      NaN

      ...      clean_essay \
0      ...      montessori school deep history student centere...

      clean_project_title  price quantity negative_sentiment_score \
0  fitness frenzy  613.96      15      0.0

      neutral_sentiment_score positive_sentiment_score compound_sentiment_score \
0      0.71      0.29      0.9829

      num_words_essay  num_words_project_title
0      99      2

[1 rows x 25 columns]
```

Quantity

```
In [117]: from sklearn.preprocessing import Normalizer
normalizerTr = Normalizer()
# normalizer.fit(X_train['price'].values)
# this will rise an error Expected 2D array, got 1D array instead:
# array=[105.22 215.96 96.01 ... 368.98 80.53 709.67].
# Reshape your data either using
# array.reshape(-1, 1) if your data has a single feature
# array.reshape(1, -1) if it contains a single sample.
normalizerTr.fit(X_train['quantity'].values.reshape(-1,1))

X_train_q_norm = normalizerTr.transform(X_train['quantity'].values.reshape(-1,1))
X_test_q_norm = normalizerTr.transform(X_test['quantity'].values.reshape(-1,1))

print("After vectorizations")
print(X_train_q_norm.shape, y_train.shape)
print(X_test_q_norm.shape, y_test.shape)
print("=="*100)
```

After vectorizations

(33500, 1) (33500,)

(16500, 1) (16500,)

=====

#Words in clean_project_title

```
In [118]: from sklearn.preprocessing import Normalizer
normalizerTr = Normalizer()
# normalizer.fit(X_train['price'].values)
# this will rise an error Expected 2D array, got 1D array instead:
# array=[105.22 215.96 96.01 ... 368.98 80.53 709.67].
# Reshape your data either using
# array.reshape(-1, 1) if your data has a single feature
# array.reshape(1, -1) if it contains a single sample.
normalizerTr.fit(X_train['num_words_project_title'].values.reshape(-1,1))

X_train_nwpt_norm = normalizerTr.transform(X_train['num_words_project_title'].values.reshape(-1,1))
X_test_nwpt_norm = normalizerTr.transform(X_test['num_words_project_title'].values.reshape(-1,1))

print("After vectorizations")
print(X_train_nwpt_norm.shape, y_train.shape)
print(X_test_nwpt_norm.shape, y_test.shape)
print("=="*100)
```

After vectorizations

(33500, 1) (33500,)

(16500, 1) (16500,)

#Words in clean_essay

```
In [119]: from sklearn.preprocessing import Normalizer
normalizerTr = Normalizer()
# normalizer.fit(X_train['price'].values)
# this will rise an error Expected 2D array, got 1D array instead:
# array=[105.22 215.96 96.01 ... 368.98 80.53 709.67].
# Reshape your data either using
# array.reshape(-1, 1) if your data has a single feature
# array.reshape(1, -1) if it contains a single sample.
normalizerTr.fit(X_train['num_words_essay'].values.reshape(-1,1))

X_train_nwe_norm = normalizerTr.transform(X_train['num_words_essay'].values.reshape(-1,1))
X_test_nwe_norm = normalizerTr.transform(X_test['num_words_essay'].values.reshape(-1,1))

print("After vectorizations")
print(X_train_nwe_norm.shape, y_train.shape)
print(X_test_nwe_norm.shape, y_test.shape)
print("=="*100)
```

After vectorizations
(33500, 1) (33500,)
(16500, 1) (16500,)

negative_sentiment_score

```
In [120]: from sklearn.preprocessing import Normalizer
normalizerTr = Normalizer()
# normalizer.fit(X_train['price'].values)
# this will rise an error Expected 2D array, got 1D array instead:
# array=[105.22 215.96 96.01 ... 368.98 80.53 709.67].
# Reshape your data either using
# array.reshape(-1, 1) if your data has a single feature
# array.reshape(1, -1) if it contains a single sample.
normalizerTr.fit(X_train['negative_sentiment_score'].values.reshape(-1,1))

X_train_nss_norm = normalizerTr.transform(X_train['negative_sentiment_score'].values.reshape(-1,1))
X_test_nss_norm = normalizerTr.transform(X_test['negative_sentiment_score'].values.reshape(-1,1))

print("After vectorizations")
print(X_train_nss_norm.shape, y_train.shape)
print(X_test_nss_norm.shape, y_test.shape)
print("=="*100)
```

After vectorizations
(33500, 1) (33500,)
(16500, 1) (16500,)

=====

neutral_sentiment_score

In [121]: `from sklearn.preprocessing import Normalizer`

```
normalizerTr = Normalizer()
# normalizer.fit(X_train['price'].values)
# this will rise an error Expected 2D array, got 1D array instead:
# array=[105.22 215.96 96.01 ... 368.98 80.53 709.67].
# Reshape your data either using
# array.reshape(-1, 1) if your data has a single feature
# array.reshape(1, -1) if it contains a single sample.
normalizerTr.fit(X_train['neutral_sentiment_score'].values.reshape(-1,1))
```

```
X_train_ness_norm = normalizerTr.transform(X_train['neutral_sentiment_score'].values.reshape(-1,1))
X_test_ness_norm = normalizerTr.transform(X_test['neutral_sentiment_score'].values.reshape(-1,1))
```

```
print("After vectorizations")
print(X_train_ness_norm.shape, y_train.shape)
print(X_test_ness_norm.shape, y_test.shape)
print("="*100)
```

After vectorizations
(33500, 1) (33500,)
(16500, 1) (16500,)

=====

positive_sentiment_score

In [122]: `from sklearn.preprocessing import Normalizer`

```
normalizerTr = Normalizer()
# normalizer.fit(X_train['price'].values)
# this will rise an error Expected 2D array, got 1D array instead:
# array=[105.22 215.96 96.01 ... 368.98 80.53 709.67].
# Reshape your data either using
# array.reshape(-1, 1) if your data has a single feature
# array.reshape(1, -1) if it contains a single sample.
normalizerTr.fit(X_train['positive_sentiment_score'].values.reshape(-1,1))
```

```
X_train_pss_norm = normalizerTr.transform(X_train['positive_sentiment_score'].values.reshape(-1,1))
X_test_pss_norm = normalizerTr.transform(X_test['positive_sentiment_score'].values.reshape(-1,1))
```

```
print("After vectorizations")
```

```

print(X_train_pss_norm.shape, y_train.shape)
print(X_test_pss_norm.shape, y_test.shape)
print("="*100)

```

After vectorizations

```

(33500, 1) (33500,)
(16500, 1) (16500,)

```

compound_sentiment_score

In [123]: `from sklearn.preprocessing import Normalizer`

```

normalizerTr = Normalizer()
# normalizer.fit(X_train['price'].values)
# this will rise an error Expected 2D array, got 1D array instead:
# array=[105.22 215.96 96.01 ... 368.98 80.53 709.67].
# Reshape your data either using
# array.reshape(-1, 1) if your data has a single feature
# array.reshape(1, -1) if it contains a single sample.
normalizerTr.fit(X_train['compound_sentiment_score'].values.reshape(-1,1))

```

```

X_train_css_norm = normalizerTr.transform(X_train['compound_sentiment_score'].values.reshape(-1,1))
X_test_css_norm = normalizerTr.transform(X_test['compound_sentiment_score'].values.reshape(-1,1))

```

```

print("After vectorizations")
print(X_train_css_norm.shape, y_train.shape)
print(X_test_css_norm.shape, y_test.shape)
print("="*100)

```

After vectorizations

```

(33500, 1) (33500,)
(16500, 1) (16500,)

```

2.2 Merging all features with newly added features using hstack

In [124]: `# merge two sparse matrices: https://stackoverflow.com/a/19710648/4084039`

```

from scipy.sparse import hstack
X_tr_new = hstack((X_train_cc_ohe, X_train_csc_ohe, X_train_grade_ohe, X_train_state_ohe, X_train_zip_ohe))
X_te_new = hstack((X_test_cc_ohe, X_test_csc_ohe, X_test_grade_ohe, X_test_state_ohe, X_test_zip_ohe))

print("Final Data matrix")
print(X_tr_new.shape, y_train.shape)
print(X_te_new.shape, y_test.shape)
print("="*100)

```


Final Data matrix
(33500, 107) (33500,)
(16500, 107) (16500,)

=====

2.3 Learning Curve (AUC)

```
In [129]: # https://scikit-learn.org/stable/modules/generated/sklearn.model\_selection.GridSearchCV.html
from sklearn.model_selection import GridSearchCV
from sklearn.linear_model import LogisticRegression

lr1 = LogisticRegression()

parameters = {'C': [10**-4, 10**-3, 10**-2, 10**-1, 10**0, 10**1, 10**2, 10**3, 10**4]}
#parameters = {'alpha':[0.0001, 0.0025, 0.0005, 0.0075, 0.001, 0.025, .005, 0.075, 0.1, 0.25, 0.5, 0.75, 1]}

clf = GridSearchCV(lr1, parameters, cv=7, scoring='roc_auc')
clr4 = clf.fit(X_tr_new, y_train)

train_auc4= clf.cv_results_['mean_train_score']
train_auc_std4= clf.cv_results_['std_train_score']
cv_auc4 = clf.cv_results_['mean_test_score']
cv_auc_std4 = clf.cv_results_['std_test_score']

plt.plot(parameters['C'], train_auc4, label='Train AUC')
# this code is copied from here: https://stackoverflow.com/a/48803361/4084039
plt.gca().fill_between(parameters['C'],train_auc4 - train_auc_std4,train_auc4 + train_auc_std4,alpha=

plt.plot(parameters['C'], cv_auc4, label='CV AUC')
# this code is copied from here: https://stackoverflow.com/a/48803361/4084039
plt.gca().fill_between(parameters['C'],cv_auc4 - cv_auc_std4, cv_auc4 + cv_auc_std4,alpha=0.2,color=

plt.scatter(parameters['C'], train_auc4, label='Train AUC points')
plt.scatter(parameters['C'], cv_auc4, label='CV AUC points')

plt.xscale('log')
plt.legend()
plt.xlabel("K: hyperparameter")
plt.ylabel("AUC")
plt.title("ERROR PLOTS")
plt.grid()
plt.show()

print(train_auc4)
print(cv_auc4)      #k_best = 10**-2
```



```
[0.5383002 0.55890896 0.5744245 0.57984447 0.58048341 0.58052433
 0.58052932 0.58053171 0.58053301]
[0.53727092 0.55397684 0.56215529 0.56096151 0.55935956 0.55882988
 0.55878161 0.55876429 0.55878504]
```

2.3.1 ROC Curve

In [140]: # https://scikit-learn.org/stable/modules/generated/sklearn.metrics.roc_curve.html#sklearn.metrics.roc_curve

```
from sklearn.metrics import roc_curve, auc

clf = LogisticRegression(C=10, penalty='l1');
clf.fit(X_tr4, y_train)
# roc_auc_score(y_true, y_score) the 2nd parameter should be probability estimates of the positive class
# not the predicted outputs

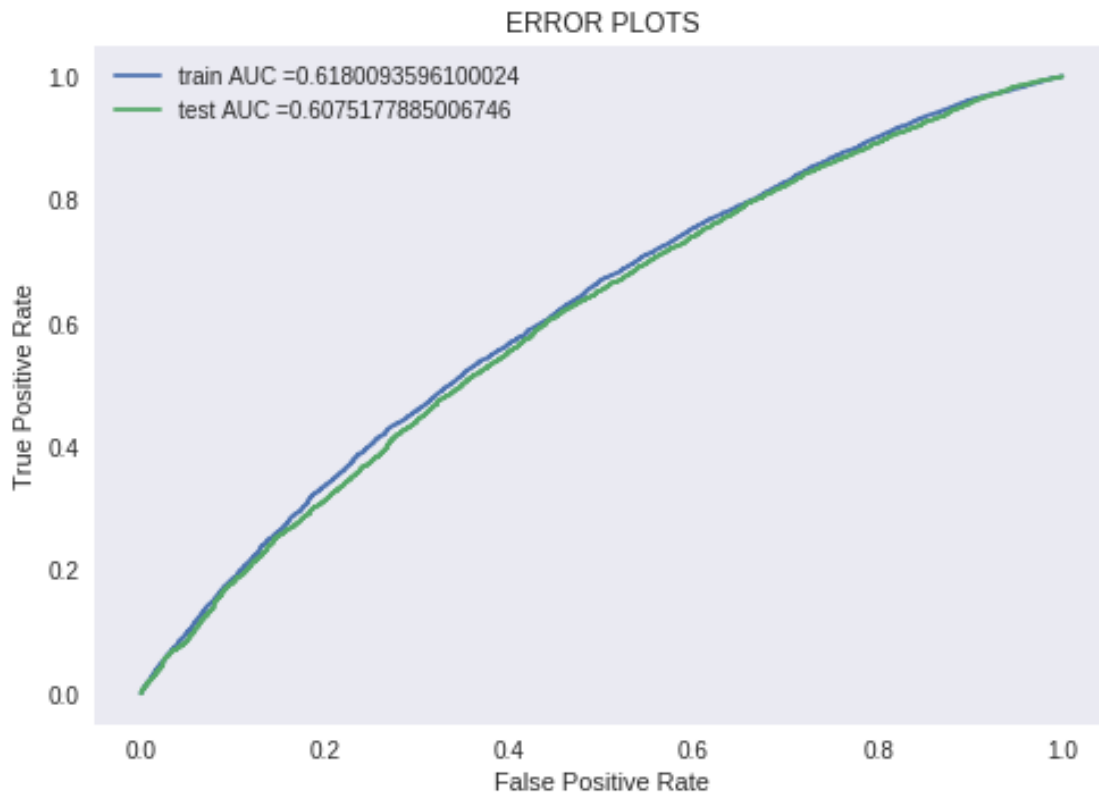
y_train_pred4 = batch_predict(clf, X_tr4)
y_test_pred4 = batch_predict(clf, X_te4)
```

```

train_fpr4, train_tpr4, tr_thresholds4 = roc_curve(y_train, y_train_pred4)
test_fpr4, test_tpr4, te_thresholds4 = roc_curve(y_test, y_test_pred4)

plt.plot(train_fpr4, train_tpr4, label="train AUC =" + str(auc(train_fpr4, train_tpr4)))
plt.plot(test_fpr4, test_tpr4, label="test AUC =" + str(auc(test_fpr4, test_tpr4)))
plt.legend()
plt.xlabel("False Positive Rate")
plt.ylabel("True Positive Rate")
plt.title("ERROR PLOTS")
plt.grid()
plt.show()

```



2.3.2 Confusion Matrix

In [0]: `def predict(proba, threshold, fpr, tpr):`

```

t = threshold[np.argmax(fpr*(1-tpr))]

```

```

# (tpr*(1-fpr)) will be maximum if your fpr is very low and tpr is very high

```

```

print("the maximum value of tpr*(1-fpr)", max(tpr*(1-fpr)), "for threshold", np.round(t,3))
predictions = []

```

```

for i in proba:
    if i>=t:
        predictions.append(1)
    else:
        predictions.append(0)
return predictions

```

In [142]: [#https://www.quantinsti.com/blog/creating-heatmap-using-python-seaborn](https://www.quantinsti.com/blog/creating-heatmap-using-python-seaborn)
import seaborn as sns; sns.set()

```

con_m_train = confusion_matrix(y_train, predict(y_train_pred4, tr_thresholds4, train_fpr4, train_tpr4))
con_m_test = confusion_matrix(y_test, predict(y_test_pred4, te_thresholds4, test_fpr4, test_tpr4))

```

```
key = (np.asarray([['TN','FP'], ['FN', 'TP']]))
```

```
fig, ax = plt.subplots(1,2, figsize=(12,5))
```

```

labels_train = (np.asarray(["{0} = {1:.2f}" .format(key, value) for key, value in zip(key.flatten(), con_m_train.flatten())])
labels_test = (np.asarray(["{0} = {1:.2f}" .format(key, value) for key, value in zip(key.flatten(), con_m_test.flatten())])

```

```

sns.heatmap(con_m_train, linewidths=.5, xticklabels=['PREDICTED : NO', 'PREDICTED : YES'], yticklabels=['ACTUAL : NO', 'ACTUAL : YES'])
sns.heatmap(con_m_test, linewidths=.5, xticklabels=['PREDICTED : NO', 'PREDICTED : YES'], yticklabels=['ACTUAL : NO', 'ACTUAL : YES'])

```

```


ax[0].set_title('Train Set')
ax[1].set_title('Test Set')

```

```
plt.show()
```

the maximum value of $tpr*(1-fpr)$ 0.3414943415840221 for threshold 0.844

the maximum value of $tpr*(1-fpr)$ 0.3377841974545248 for threshold 0.846



3. Conclusions

In [143]: # Please compare all your models using Prettytable library

```
# Please compare all your models using Prettytable library
```

```
from prettytable import PrettyTable
```

```
x = PrettyTable()
```

```
x.field_names = ["Vectorizer", "Model", "HyperParameter", "AUC"]
```

```
x.add_row(["BOW", "Logistic Regression", 10**-2, 0.592])
```

```
x.add_row(["TFIDF", "Logistic Regression", 1.0, 0.655])
```

```
x.add_row(["Avg W2V", "Logistic Regression", 1.0, 0.588])
```

```
x.add_row(["TFIDF-W2v", "Logistic Regression", 10**4, 0.608])
```

```
x.add_row(["SET5", "Logistic Regression", 10, 0.607])
```

```
print(x)
```

```
+-----+-----+-----+-----+
| Vectorizer | Model | HyperParameter | AUC |
+-----+-----+-----+-----+
| BOW | Logistic Regression | 0.01 | 0.592 |
| TFIDF | Logistic Regression | 1.0 | 0.655 |
| Avg W2V | Logistic Regression | 1.0 | 0.588 |
| TFIDF-W2v | Logistic Regression | 10000 | 0.608 |
| SET5 | Logistic Regression | 10 | 0.607 |
+-----+-----+-----+-----+
```

2.3.3 Observations

As it can be seen from the above table, that the model is performing better than random model, from all the sets, TFIDF is working fairly well having AUC score of 0.66

2.3.4 Conclusions

I took 50000 datapoints for my analysis and building my model

- I splitted the dataset into train, cv and test dataset
- Preprocessed all the text fetatures
- Vectorized all the text, categorical and numerical features, for text i used BOW & TFIDF
- Merged all features using hstack as instructed
- Using train dataset, i plotted my AUC curve using GridSearchCV using 3Fold Cross Validation for both categories
- from AUC curve, i picked best alpha. using best alpha, i plotted ROC curve on train and test data.

- Then i plotted my confusion matrix for both the sets.
- Atlast you can see my result in tabular format.